

Cardano.BM - benchmarking and logging

Alexander Diemand

Andreas Triantafyllos

November 2018

Abstract

This is a framework that combines logging, benchmarking and monitoring. Complex evaluations of STM or monadic actions can be observed from outside while reading operating system counters before and after, and calculating their differences, thus relating resource usage to such actions. Through interactive configuration, the runtime behaviour of logging or the measurement of resource usage can be altered. Further reduction in logging can be achieved by redirecting log messages to an aggregation function which will output the running statistics with less frequency than the original message.

Contents

1	Cardano BM	3
1.1	Overview	3
1.2	Introduction	3
1.2.1	Logging with Trace	3
1.2.2	Setup procedure	3
1.2.3	Measuring <i>Observables</i>	3
1.2.4	Information reduction in Aggregation	3
1.2.5	Output selection	3
1.2.6	Monitoring	3
1.3	Examples	3
1.3.1	Observing evaluation of a STM action	3
1.3.2	Observing evaluation of a monad action	3
1.3.3	Simple example showing plain logging	3
1.3.4	Complex example showing logging, aggregation of log items, and observing <i>IO</i> actions	5
1.4	Code listings	8
1.4.1	Cardano.BM.Observer.STM	8
1.4.2	Cardano.BM.Observer.Monad	9
1.4.3	BaseTrace	12
1.4.4	Cardano.BM.Trace	13
1.4.5	Cardano.BM.Setup	18
1.4.6	Cardano.BM.Counters	20
1.4.7	Cardano.BM.Counters.Common	21
1.4.8	Cardano.BM.Counters.Dummy	22
1.4.9	Cardano.BM.Counters.Linux	22
1.4.10	Cardano.BM.Data.Aggregated	28
1.4.11	Cardano.BM.Data.Backend	32
1.4.12	Cardano.BM.Data.Configuration	33
1.4.13	Cardano.BM.Data.Counter	34
1.4.14	Cardano.BM.Data.LogItem	36
1.4.15	Cardano.BM.Data.Observable	37
1.4.16	Cardano.BM.Data.Output	37
1.4.17	Cardano.BM.Data.Severity	38
1.4.18	Cardano.BM.Data.SubTrace	38
1.4.19	Cardano.BM.Data.Trace	39
1.4.20	Cardano.BM.Configuration	39
1.4.21	Cardano.BM.Configuration.Model	40
1.4.22	Cardano.BM.Output.Switchboard	46

1.4.23	Cardano.BM.Output.Log	49
1.4.24	Cardano.BM.Output.EKGView	55
1.4.25	Cardano.BM.Output.Aggregation	58

Chapter 1

Cardano BM

1.1 Overview

In figure 1.1 we display the relationships among modules in *Cardano.BM*. The arrows indicate import of a module. The arrows with a triangle at one end would signify "inheritance" in object-oriented programming, but we use it to show that one module replaces the other in the namespace, thus refines its interface.

1.2 Introduction

1.2.1 Logging with **Trace**

1.2.2 Setup procedure

Hierarchy of **Traces**

1.2.3 Measuring *Observables*

1.2.4 Information reduction in **Aggregation**

1.2.5 Output selection

1.2.6 Monitoring

1.3 Examples

1.3.1 Observing evaluation of a STM action

1.3.2 Observing evaluation of a monad action

1.3.3 Simple example showing plain logging

```
{-# LANGUAGE OverloadedStrings #-}
module Main
  (main)
  where
import Control.Concurrent (threadDelay)
import Cardano.BM.Configuration.Static (defaultConfigStdout)
```

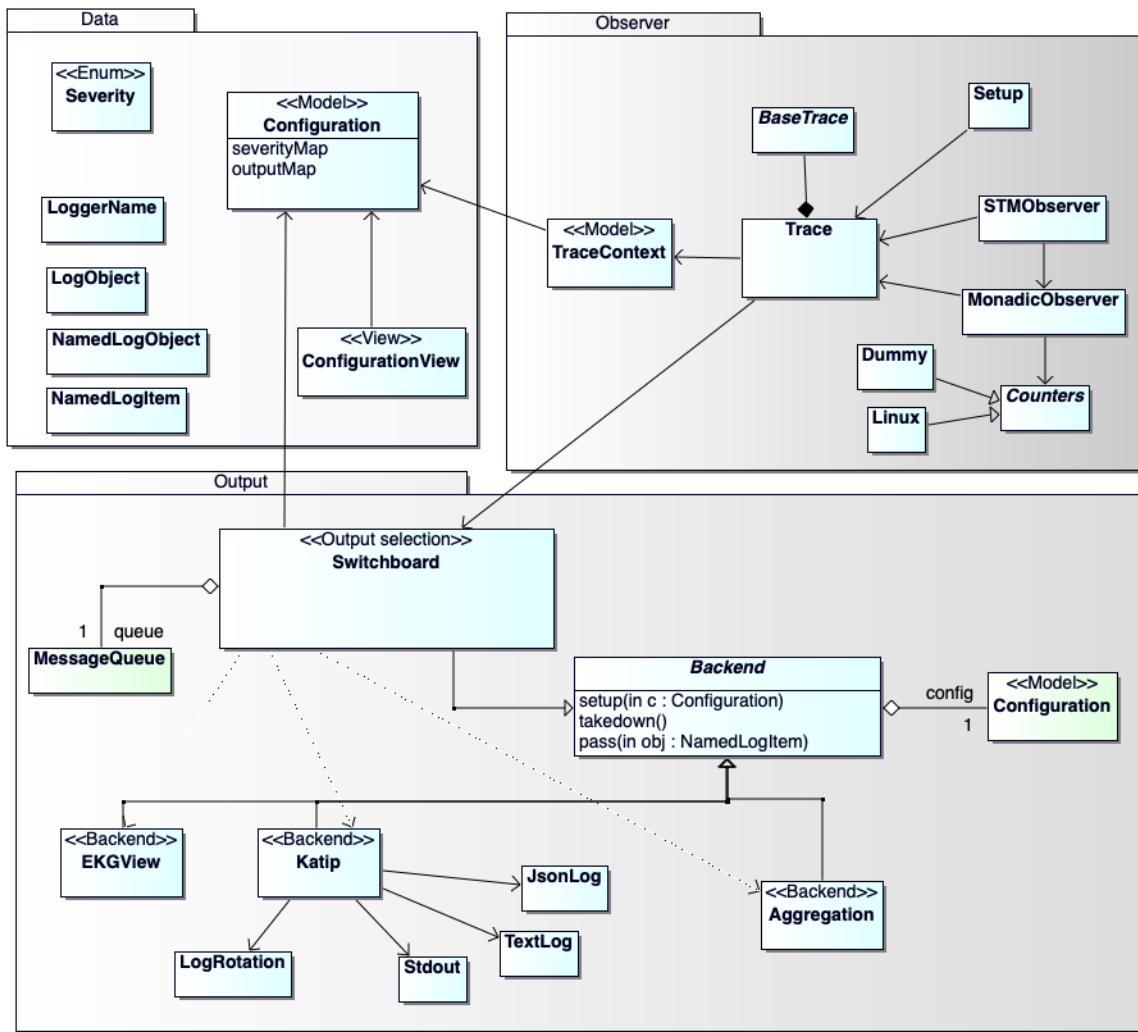


Figure 1.1: Overview of module relationships

```

import Cardano.BM.Setup (setupTrace)
import Cardano.BM.Trace (logDebug, logError, logInfo, logNotice,
                        logWarning)

main :: IO ()
main = do
  c ← defaultConfigStdout
  tr ← setupTrace (Right c) "simple"
  logDebug tr "this is a debug message"
  logInfo tr "this is an information."
  logNotice tr "this is a notice!"
  logWarning tr "this is a warning!"
  logError tr "this is an error!"
  threadDelay 80000

```



Figure 1.2: Setup procedure

```
return ()
```

1.3.4 Complex example showing logging, aggregation of log items, and observing IO actions

Module header and import directives

```

{-# LANGUAGE OverloadedStrings #-}
module Main
  (main)
  where
import qualified Control.Concurrent.Async as Async
import Control.Concurrent (threadDelay)
import qualified Cardano.BM.Configuration.Model as CM
import Cardano.BM.Data.Aggregated (Measurable (..))
import Cardano.BM.Data.BackendKind
import Cardano.BM.Data.LogItem
import Cardano.BM.Data.Observable
import Cardano.BM.Data.Output
import Cardano.BM.Data.Severity
import Cardano.BM.Data.SubTrace
import Cardano.BM.Observer.Monad (bracketObserveIO)
import Cardano.BM.Setup
import Cardano.BM.Trace
import System.Random
  
```

Define configuration

The output can be viewed in EKG on <http://localhost:12789>.

```

config :: IO CM.Configuration
config = do
  c ← CM.empty
  CM.setMinSeverity c Debug
  CM.setSetupBackends c [KatipBK, AggregationBK, EKGViewBK]
  {-per default each messages is sent to the logs, if not otherwise defined (see below: 'CM.setBackend') -}
  CM.setDefaultBackends c [KatipBK]
  CM.setSetupScribes c [ScribeDefinition {
    scName = "stdout"
    ,scKind = StdoutSK
    ,scRotation = Nothing
  }
  ,ScribeDefinition {
    scName = "out.json"
    ,scKind = FileJsonSK
    ,scRotation = Nothing
  }
  ,ScribeDefinition {
    scName = "out.txt"
    ,scKind = FileTextSK
    ,scRotation = Nothing
  }
  ]
  -- per default each messages is sent to the logs, if not otherwise defined (see below:
  CM.setDefaultScribes c ["StdoutSK::stdout", "FileJsonSK::out.json"]
  CM.setScribes c "complex.random" (Just ["StdoutSK::stdout", "FileTextSK::out.txt"])
  CM.setScribes c "complex.random.aggregated" (Just ["StdoutSK::stdout"])
  {-define a subtrace whose behaviour is to copy all log items, and pass them up with a name added to th
  CM.setSubTrace c "complex.random" (Just $ TeeTrace "copy")
  -- not every aggregated value needs to be displayed
  CM.setSubTrace c "#ekgview"
    (Just $ FilterTrace [Drop (StartsWith "#ekgview.#aggregation.complex.random"),
      Unhide (Named "count"),
      Unhide (Named "mean")
    ])
  {-define a subtrace whose behaviour is to copy all log items, and pass them up with a name added to th
  CM.setSubTrace c "complex.observeIO" (Just $ ObservableTrace [GhcRtsStats, MemoryStats])
  -- forward the random number to aggregation:
  CM.setBackends c "complex.random" (Just [AggregationBK, KatipBK])
  CM.setBackends c "complex.random.copy" (Just [AggregationBK])
  -- forward the observed values to aggregation:
  CM.setBackends c "complex.observeIO" (Just [KatipBK])
  -- forward the aggregated output to the EKG view:
  CM.setBackends c "#aggregation.complex.random" (Just [EKGViewBK])
  CM.setBackends c "#aggregation.complex.random.copy" (Just [EKGViewBK])
  CM.setBackends c "#aggregation.complex.observeIO" (Just [EKGViewBK])
  -- start EKG on http://localhost:12789

```



```
CM.setEKGport c 12789
return c
```

Thread that outputs a random number to a **Trace**

```
randomThr :: Trace IO → IO (Async.Async ())
randomThr trace = do
  logInfo trace "starting random generator"
  trace' ← subTrace "random" trace
  proc ← Async.async (loop trace')
  return proc
where
  loop tr = do
    threadDelay 800000
    num ← randomRIO (42 - 42, 42 + 42) :: IO Double
    traceNamedObject tr (LP (LogValue "rr" (PureD num)))
    loop tr
```

Thread that observes an IO action

```
observeIO :: Trace IO → IO (Async.Async ())
observeIO trace = do
  logInfo trace "starting observer"
  proc ← Async.async (loop trace)
  return proc
where
  loop tr = do
    threadDelay 1000000 -- 1 second
    bracketObserveIO trace "observeIO" $ do
      num ← randomRIO (100000, 2000000) :: IO Int
      _ ← return $ reverse $ reverse $ 42 : [1..num]
      threadDelay 50000 -- .05 second
      pure ()
    loop tr
```

Main entry point

```
main :: IO ()
main = do
  -- create configuration
  c ← config
  -- create initial top-level Trace
  tr ← setupTrace (Right c) "complex"
  logNotice tr "starting program; hit CTRL-C to terminate"
```

```

logInfo tr "watch its progress on http://localhost:12789"
  {-start thread sending unbounded sequence of random numbers to a trace which aggregates them into a
  procRandom ← randomThr tr
  -- start thread endlessly reversing lists of random length
  procObsvIO ← observeIO tr
  -- wait for observer thread to finish, ignoring any exception
  _ ← Async.waitCatch procObsvIO
  -- wait for random thread to finish, ignoring any exception
  _ ← Async.waitCatch procRandom
  return ()

```

1.4 Code listings

1.4.1 Cardano.BM.Observer.STM

```

stmWithLog :: STM.STM (t,[LogObject]) → STM.STM (t,[LogObject])
stmWithLog action = action

```

Observe STM action in a named context

With given name, create a **SubTrace** according to **Configuration** and run the passed STM action on it.

```

bracketObserveIO :: Trace IO → Text → STM.STM t → IO t
bracketObserveIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  let subtrace = typeofTrace logTrace
  bracketObserveIO' subtrace logTrace action
where
  bracketObserveIO' :: SubTrace → Trace IO → STM.STM t → IO t
  bracketObserveIO' NoTrace _ act =
    STM.atomically act
  bracketObserveIO' subtrace logTrace act = do
    mCountersid ← observeOpen subtrace logTrace
    -- run action; if an exception is caught will be logged and rethrown.
    t ← (STM.atomically act) 'catch' (λ(e :: SomeException) → (logError logTrace (pack (show e)) >> throwM e))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        logNotice logTrace ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace logTrace countersid [ ]
        case res of
          Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))

```

```

    _ → pure ()
  pure t

```

Observe STM action in a named context and output captured log items

The STM action might output messages, which after "success" will be forwarded to the logging trace. Otherwise, this function behaves the same as **Observe STM action in a named context**.

```

bracketObserveLogIO :: Trace IO → Text → STM.STM (t, [LogObject]) → IO t
bracketObserveLogIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  let subtrace = typeofTrace logTrace
  bracketObserveLogIO' subtrace logTrace action
where
  bracketObserveLogIO' :: SubTrace → Trace IO → STM.STM (t, [LogObject]) → IO t
  bracketObserveLogIO' NoTrace _ act = do
    (t, _) ← STM.atomically $ stmWithLog act
    pure t
  bracketObserveLogIO' subtrace logTrace act = do
    mCountersid ← observeOpen subtrace logTrace
    -- run action, return result and log items; if an exception is
    -- caught will be logged and rethrown.
    (t, as) ← (STM.atomically $ stmWithLog act) 'catch'
      (λ(e :: SomeException) → (logError logTrace (pack (show e)) >> throwM e))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        logNotice logTrace ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace logTrace countersid as
        case res of
          Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
          _ → pure ()
    pure t

```

1.4.2 Cardano.BM.Observer.Monad

Monadic.bracketObserverIO

Observes an IO action and adds a name to the logger name of the passed in **Trace**. An empty Text leaves the logger name untouched.

Microbenchmarking steps:

1. Create a trace which will have been configured to observe things besides logging.

```

import qualified Cardano.BM.Configuration.Model as CM
ooo
c ← config
trace@(ctx, _) ← setupTrace (Right c) "demo-playground"
where
  config :: IO CM.Configuration
  config = do
    c ← CM.empty
    CM.setMinSeverity c Debug
    CM.setSetupBackends c [KatipBK, AggregationBK]
    CM.setDefaultBackends c [KatipBK, AggregationBK]
    CM.setSetupScribes c [ScribeDefinition {
      scName = "stdout"
    , scKind = StdoutSK
    , scRotation = Nothing
    }
    ]
    CM.setDefaultScribes c ["StdoutSK::stdout"]
  return c

```

2. *c* is the **Configuration** of *trace*. In order to enable the collection and processing of measurements (min, max, mean, std-dev) *AggregationBK* is needed.

```
CM.setDefaultBackends c [KatipBK, AggregationBK]
```

in a configuration file (YAML) means

```

defaultBackends:
  - KatipBK
  - AggregationBK

```

3. Set the measurements that you want to take by changing the configuration of the *trace* using *setSubTrace*, in order to declare the namespace where we want to enable the particular measurements and the list with the kind of measurements.

```

CM.setSubTrace
  (configuration ctx)
  "demo-playground.submit-tx"
  (Just $ ObservableTrace observablesSet)
where
  observablesSet = [MonotonicClock, MemoryStats]

```

4. Find an action to measure. e.g.:

```
runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())
```

and use **bracketObserveIO**. e.g.:

```

bracketObserveIO trace "submit-tx" $
  runProtocolWithPipe x hdl proto 'catch' (λProtocolStopped → return ())

```

```

bracketObserveIO :: Trace IO → Text → IO t → IO t
bracketObserveIO logTrace0 name action = do
  logTrace ← subTrace name logTrace0
  bracketObserveIO' (typeofTrace logTrace) logTrace action
where
  bracketObserveIO' :: SubTrace → Trace IO → IO t → IO t
  bracketObserveIO' NoTrace _ act = act
  bracketObserveIO' subtrace logTrace act = do
    mCountersid ← observeOpen subtrace logTrace
    -- run action; if an exception is caught will be logged and rethrown.
    t ← act 'catch' (λ(e :: SomeException) → (logError logTrace (pack (show e)) >> throwM e))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        logNotice logTrace ("ObserveOpen: " <> pack (show openException))
      Right countersid → do
        res ← observeClose subtrace logTrace countersid [ ]
        case res of
          Left ex → logNotice logTrace ("ObserveClose: " <> pack (show ex))
          _ → pure ()
    pure t

```

Monadic.bracketObserverM

Observes a *MonadIO* $m \Rightarrow m$ action and adds a name to the logger name of the passed in **Trace**. An empty *Text* leaves the logger name untouched.

```

bracketObserveM :: (MonadCatch m, MonadIO m) ⇒ Trace IO → Text → m t → m t
bracketObserveM logTrace0 name action = do
  logTrace ← liftIO $ subTrace name logTrace0
  bracketObserveM' (typeofTrace logTrace) logTrace action
where
  bracketObserveM' :: (MonadCatch m, MonadIO m) ⇒ SubTrace → Trace IO → m t → m t
  bracketObserveM' NoTrace _ act = act
  bracketObserveM' subtrace logTrace act = do
    mCountersid ← liftIO $ observeOpen subtrace logTrace
    -- run action; if an exception is caught will be logged and rethrown.
    t ← act 'catch'
      (λ(e :: SomeException) → (liftIO (logError logTrace (pack (show e)) >> throwM e)))
    case mCountersid of
      Left openException →
        -- since observeOpen faced an exception there is no reason to call observeClose
        -- however the result of the action is returned
        liftIO $ logNotice logTrace ("ObserveOpen: " <> pack (show openException))

```

```

Right countersid → do
  res ← liftIO $ observeClose subtrace logTrace countersid [ ]
  case res of
    Left ex → liftIO (logNotice logTrace ("ObserveClose: " <> pack (show ex)))
    _ → pure ()
pure t

```

observerOpen

```

observeOpen :: SubTrace → Trace IO → IO (Either SomeException CounterState)
observeOpen subtrace logTrace = (do
  identifier ← newUnique
  -- take measurement
  counters ← readCounters subtrace
  let state = CounterState identifier counters
  -- send opening message to Trace
  traceNamedObject logTrace $ ObserveOpen state
  return (Right state)) 'catch' (return ∘ Left)

```

observeClose

```

observeClose :: SubTrace → Trace IO → CounterState → [LogObject] → IO (Either SomeException ())
observeClose subtrace logTrace initState logObjects = (do
  let identifier = csIdentifier initState
  initialCounters = csCounters initState
  -- take measurement
  counters ← readCounters subtrace
  -- send closing message to Trace
  traceNamedObject logTrace $ ObserveClose (CounterState identifier counters)
  -- send diff message to Trace
  traceNamedObject logTrace $
    ObserveDiff (CounterState identifier (diffCounters initialCounters counters))
  -- trace the messages gathered from inside the action
  forM_ logObjects $ traceNamedObject logTrace
  return (Right ())) 'catch' (return ∘ Left)

```

1.4.3 BaseTrace

Contravariant

A covariant is a functor: $F A \rightarrow F B$

A contravariant is a functor: $F B \rightarrow F A$

$Op\ a\ b$ implements the inverse to 'arrow' $getOp :: b \rightarrow a$, which when applied to a **BaseTrace** of type $Op\ (m\ ())\ s$, yields $s \rightarrow m\ ()$. In our case, Op accepts an action in a monad m with input type **LogNamed LogObject** (see 'Trace').

```
newtype BaseTrace m s = BaseTrace {runTrace :: Op (m ()) s}
```

contramap

A covariant functor defines the function " $fmap :: (a \rightarrow b) \rightarrow f\ a \rightarrow f\ b$ ". In case of a contravariant functor, it is the dual function " $contramap :: (a \rightarrow b) \rightarrow f\ b \rightarrow f\ a$ " which is defined.

In the following instance, `runTrace` extracts type " $Op\ (m\ ())\ s$ " to which `contramap` applies f , thus " $f\ s \rightarrow m\ ()$ ". The constructor `BaseTrace` restores " $Op\ (m\ ())\ (f\ s)$ ".

```
instance Contravariant (BaseTrace m) where
  contramap f = BaseTrace  $\circ$  contramap f  $\circ$  runTrace
```

traceWith

Accepts a `Trace` and some payload s . First it gets the contravariant from the `Trace` as type " $Op\ (m\ ())\ s$ " and, after " $getOp :: b \rightarrow a$ " which translates to " $s \rightarrow m\ ()$ ", calls the action on the `LogNamed LogObject`.

```
traceWith :: BaseTrace m s  $\rightarrow s \rightarrow m\ ()$ 
traceWith = getOp  $\circ$  runTrace
```

natTrace

Natural transformation from monad m to monad n .

```
natTrace :: (forall x  $\circ m\ x \rightarrow n\ x$ )  $\rightarrow$  BaseTrace m s  $\rightarrow$  BaseTrace n s
natTrace nat (BaseTrace (Op tr)) = BaseTrace $ Op $ nat  $\circ$  tr
```

noTrace

A `Trace` that discards all inputs.

```
noTrace :: Applicative m  $\Rightarrow$  BaseTrace m a
noTrace = BaseTrace $ Op $ const (pure ())
```

1.4.4 Cardano.BM.Trace

Utilities

Natural transformation from monad m to monad n .

```
natTrace :: (forall x  $\circ m\ x \rightarrow n\ x$ )  $\rightarrow$  Trace m  $\rightarrow$  Trace n
natTrace nat (ctx, trace) = (ctx, BaseTrace.natTrace nat trace)
```

Access type of `Trace`.

```
typeofTrace :: Trace m  $\rightarrow$  SubTrace
typeofTrace (ctx, _) = tracetype ctx
```

Update type of `Trace`.

```
updateTracetype :: SubTrace  $\rightarrow$  Trace m  $\rightarrow$  Trace m
updateTracetype subtr (ctx, tr) = (ctx { tracetype = subtr }, tr)
```

Enter new named context

The context name is created and checked that its size is below a limit (currently 80 chars). The minimum severity that a log message must be labelled with is looked up in the configuration and recalculated.

```

appendName :: MonadIO m => LoggerName -> Trace m -> m (Trace m)
appendName name (ctx, trace) = do
  let prevLoggerName = loggerName ctx
      prevMinSeverity = minSeverity ctx
      newLoggerName = appendWithDot prevLoggerName name
      globMinSeverity ← liftIO $ Config.minSeverity (configuration ctx)
      namedSeverity ← liftIO $ Config.inspectSeverity (configuration ctx) newLoggerName
  case namedSeverity of
    Nothing -> return (ctx {loggerName = newLoggerName}, trace)
    Just sev -> return (ctx {loggerName = newLoggerName
                          , minSeverity = max (max sev prevMinSeverity) globMinSeverity}
                      , trace)

appendWithDot :: LoggerName -> LoggerName -> LoggerName
appendWithDot "" newName = T.take 80 newName
appendWithDot xs "" = xs
appendWithDot xs newName = T.take 80 $ xs <> "." <> newName

```

Contramap a trace and produce the naming context

```

named :: BaseTrace.BaseTrace m (LogNamed i) -> LoggerName -> BaseTrace.BaseTrace m i
named trace name = contramap (LogNamed name) trace

```

Trace a LogObject through

```

traceNamedObject
  :: MonadIO m
  => Trace m
  -> LogObject
  -> m ()
traceNamedObject trace@(ctx, logTrace) lo = do
  let lname = loggerName ctx
  doOutput ← case (typeofTrace trace) of
    FilterTrace filters -> do
      return $ evalFilters filters lname lo
    TeeTrace secName -> do
      -- create a newly named copy of the LogObject
      BaseTrace.traceWith (named logTrace (lname <> "." <> secName)) lo
      return True
  -> return True
  if doOutput

```



```

then BaseTrace.traceWith (named logTrace lname) lo
else return ()

```

Evaluation of FilterTrace

```

evalFilters :: [NameOperator] → LoggerName → LogObject → Bool
evalFilters nos nm lo =
  any (evalFilter nm lo) nos
where
  evalFilter :: LoggerName → LogObject → NameOperator → Bool
  evalFilter name item (Drop sel) = ¬ (matchName name sel) ∧ ¬ (matchItem item sel)
  evalFilter name item (Unhide sel) = matchName name sel ∨ matchItem item sel
  matchName :: LoggerName → NameSelector → Bool
  matchName name (Exact name') = name ≡ name'
  matchName name (StartsWith prefix) = T.isPrefixOf prefix name
  matchName name (EndsWith postfix) = T.isSuffixOf postfix name
  matchName name (Contains name') = T.isInfixOf name' name
  matchName _ _ = False
  matchItem :: LogObject → NameSelector → Bool
  matchItem (LP (LogValue name _)) (Named name') = name ≡ name'
  matchItem _ _ = False

```

Concrete Trace on stdout

This function returns a trace with an action of type "(LogNamed LogObject) → IO ()" which will output a text message as text and all others as JSON encoded representation to the console.

TODO remove locallock

```

locallock :: MVar ()
locallock = unsafePerformIO $ newMVar ()

stdoutTrace :: TraceNamed IO
stdoutTrace = BaseTrace.BaseTrace $ Op $ λlognamed →
  withMVar locallock $ \_ →
    case lnItem lognamed of
      LP (LogMessage logItem) →
        output (lnName lognamed) $ liPayload logItem
      obj →
        output (lnName lognamed) $ toStrict (encodeToLazyText obj)
where
  output nm msg = TIO.putStrLn $ nm <> " :: " <> msg

```

Concrete Trace into a TVar

```

traceInTVar :: STM.TVar [a] → BaseTrace.BaseTrace STM.STM a
traceInTVar tvar = BaseTrace.BaseTrace $ Op $ λa → STM.modifyTVar tvar ((:) a)

traceInTVarIO :: STM.TVar [LogObject] → TraceNamed IO
traceInTVarIO tvar = BaseTrace.BaseTrace $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) (lnItem ln))

traceNamedInTVarIO :: STM.TVar [LogNamed LogObject] → TraceNamed IO
traceNamedInTVarIO tvar = BaseTrace.BaseTrace $ Op $ λln →
  STM.atomically $ STM.modifyTVar tvar ((:) ln)

```

Check a log item's severity against the Trace's minimum severity

do we need three different **minSeverity** defined?

We do a lookup of the global **minSeverity** in the configuration. And, a lookup of the **minSeverity** for the current named context. These values might have changed in the meanwhile. A third filter is the **minSeverity** defined in the current context.

```

traceConditionally
  :: MonadIO m
  ⇒ Trace m → LogObject
  → m ()

traceConditionally logTrace@(ctx, _) msg@(LP (LogMessage item)) = do
  globminsev ← liftIO $ Config.minSeverity (configuration ctx)
  globnamesev ← liftIO $ Config.inspectSeverity (configuration ctx) (loggerName ctx)
  let minsev = max (minSeverity ctx) $ max globminsev (fromMaybe Debug globnamesev)
  flag = (liSeverity item) ≥ minsev
  when flag $ traceNamedObject logTrace msg
traceConditionally logTrace logObject =
  traceNamedObject logTrace logObject

```

Enter message into a trace

The function **traceNamedItem** creates a **LogObject** and threads this through the action defined in the **Trace**.

```

traceNamedItem
  :: (MonadIO m)
  ⇒ Trace m
  → LogSelection
  → Severity
  → T.Text
  → m ()

traceNamedItem trace p s m =
  let logmsg = LP $ LogMessage $ LogItem {liSelection = p
    , liSeverity = s

```

```

    ,liPayload = m
  }
in
  traceConditionally trace $ logmsg

```

Logging functions

```

logDebug, logInfo, logNotice, logWarning, logError, logCritical, logAlert, logEmergency
  :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebug  logTrace = traceNamedItem logTrace Both Debug
logInfo    logTrace = traceNamedItem logTrace Both Info
logNotice  logTrace = traceNamedItem logTrace Both Notice
logWarning logTrace = traceNamedItem logTrace Both Warning
logError   logTrace = traceNamedItem logTrace Both Error
logCritical logTrace = traceNamedItem logTrace Both Critical
logAlert   logTrace = traceNamedItem logTrace Both Alert
logEmergency logTrace = traceNamedItem logTrace Both Emergency

logDebugS, logInfoS, logNoticeS, logWarningS, logErrorS, logCriticalS, logAlertS, logEmergencyS
  :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebugS  logTrace = traceNamedItem logTrace Private Debug
logInfoS    logTrace = traceNamedItem logTrace Private Info
logNoticeS  logTrace = traceNamedItem logTrace Private Notice
logWarningS logTrace = traceNamedItem logTrace Private Warning
logErrorS   logTrace = traceNamedItem logTrace Private Error
logCriticalS logTrace = traceNamedItem logTrace Private Critical
logAlertS   logTrace = traceNamedItem logTrace Private Alert
logEmergencyS logTrace = traceNamedItem logTrace Private Emergency

logDebugP, logInfoP, logNoticeP, logWarningP, logErrorP, logCriticalP, logAlertP, logEmergencyP
  :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebugP  logTrace = traceNamedItem logTrace Public Debug
logInfoP    logTrace = traceNamedItem logTrace Public Info
logNoticeP  logTrace = traceNamedItem logTrace Public Notice
logWarningP logTrace = traceNamedItem logTrace Public Warning
logErrorP   logTrace = traceNamedItem logTrace Public Error
logCriticalP logTrace = traceNamedItem logTrace Public Critical
logAlertP   logTrace = traceNamedItem logTrace Public Alert
logEmergencyP logTrace = traceNamedItem logTrace Public Emergency

logDebugUnsafeP, logInfoUnsafeP, logNoticeUnsafeP, logWarningUnsafeP, logErrorUnsafeP,
  logCriticalUnsafeP, logAlertUnsafeP, logEmergencyUnsafeP
  :: (MonadIO m) => Trace m -> T.Text -> m ()
logDebugUnsafeP  logTrace = traceNamedItem logTrace PublicUnsafe Debug
logInfoUnsafeP    logTrace = traceNamedItem logTrace PublicUnsafe Info
logNoticeUnsafeP  logTrace = traceNamedItem logTrace PublicUnsafe Notice
logWarningUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Warning
logErrorUnsafeP   logTrace = traceNamedItem logTrace PublicUnsafe Error
logCriticalUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Critical

```

```
logAlertUnsafeP    logTrace = traceNamedItem logTrace PublicUnsafe Alert
logEmergencyUnsafeP logTrace = traceNamedItem logTrace PublicUnsafe Emergency
```

subTrace

Transforms the input **Trace** according to the **Configuration** using the logger name of the current **Trace** appended with the new name. If the empty *Text* is passed, then the logger name remains untouched.

```
subTrace :: MonadIO m => T.Text -> Trace m -> m (Trace m)
subTrace name tr@(ctx, _) = do
  let newName = appendWithDot (loggerName ctx) name
  subtrace0 <- liftIO $ Config.findSubTrace (configuration ctx) newName
  let subtrace = case subtrace0 of Nothing -> Neutral; Just str -> str
  case subtrace of
    Neutral      -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    UntimedTrace -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    TeeTrace _   -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    FilterTrace _ -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
    NoTrace      -> return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $ \_ -> pure ())
    DropOpening  -> return $ updateTracetype subtrace (ctx, BaseTrace.BaseTrace $ Op $ \lognamed -> do
      case lnItem lognamed of
        ObserveOpen _ -> return ()
        obj -> traceNamedObject tr obj)
    ObservableTrace _ -> do
      tr' <- appendName name tr
      return $ updateTracetype subtrace tr'
```

1.4.5 Cardano.BM.Setup

setupTrace

Setup a new **Trace** (**Trace**) with either a given **Configuration** (**Configuration.Model**) or a *FilePath* to a configuration file.

```
setupTrace :: MonadIO m => Either FilePath Config.Configuration -> Text -> m (Trace m)
setupTrace (Left cfgFile) name = do
  c <- liftIO $ Config.setup cfgFile
  setupTrace_ c name
setupTrace (Right c) name = setupTrace_ c name
```




Figure 1.3: Setup procedure

```

setupTrace_ :: MonadIO m => Config.Configuration -> Text -> m (Trace m)
setupTrace_ c name = do
  sb <- liftIO $ Switchboard.realize c
  sev <- liftIO $ Config.minSeverity c
  ctx <- liftIO $ newContext name c sev sb
  let logTrace = natTrace liftIO (ctx, Switchboard.mainTrace sb)
  logTrace' <- subTrace " " logTrace
  return logTrace'

```

withTrace

```

withTrace :: MonadIO m => Config.Configuration -> Text -> (Trace m -> m t) -> m t
withTrace cfg name action = do
  logTrace <- setupTrace (Right cfg) name
  action logTrace

```

newContext

```

newContext :: LoggerName
  -> Config.Configuration
  -> Severity
  -> Switchboard.Switchboard
  -> IO TraceContext
newContext name cfg sev sb = do
  return $ TraceContext {
    loggerName = name
    , configuration = cfg
    , minSeverity = sev
    , tracetype = Neutral
  }

```

```
,shutdown = unrealize sb
}
```

1.4.6 Cardano.BM.Counters

Here the platform is chosen on which we compile this program.

Currently, we mainly support *Linux* with its 'proc' filesystem.

```
{-# LANGUAGE CPP #-}
# if defined (linux_HOST_OS)
# define LINUX
# endif
module Cardano.BM.Counters
(
    Platform.readCounters
    ,diffTimeObserved
    ,getMonoClock
) where
# ifdef LINUX
import qualified Cardano.BM.Counters.Linux as Platform
# else
import qualified Cardano.BM.Counters.Dummy as Platform
# endif
import Cardano.BM.Counters.Common (getMonoClock)
import Cardano.BM.Data.Aggregated (Measurable (..))
import Cardano.BM.Data.Counter
import Data.Time.Units (Microsecond)
```

Calculate difference between clocks

```
diffTimeObserved :: CounterState → CounterState → Microsecond
diffTimeObserved (CounterState id0 startCounters) (CounterState id1 endCounters) =
    let
        startTime = getMonotonicTime startCounters
        endTime = getMonotonicTime endCounters
    in
        if (id0 == id1)
        then endTime - startTime
        else error "these clocks are not from the same experiment"
    where
        getMonotonicTime counters = case (filter isMonotonicClockCounter counters) of
            [(Counter MonotonicClockTime _ (Microseconds micros))] → fromInteger micros
            _ → error "A time measurement is missing!"
        isMonotonicClockCounter :: Counter → Bool
        isMonotonicClockCounter = (MonotonicClockTime ==) ∘ cType
```

1.4.7 Cardano.BM.Counters.Common

Common functions that serve *readCounters* on all platforms.

```
nominalTimeToMicroseconds :: Word64 → Microsecond
nominalTimeToMicroseconds = fromMicroseconds ∘ toInteger ∘ ('div' 1000)
```

Read monotonic clock

```
getMonoClock :: IO [Counter]
getMonoClock = do
  t ← getMonotonicTimeNSec
  return [Counter MonotonicClockTime "monoclock" $ Microseconds (toInteger $ nominalTimeToMicroseconds t)]
```

Read GHC RTS statistics

Read counters from GHC's RTS (runtime system). The values returned are as per the last GC (garbage collection) run.

```
readRTSStats :: IO [Counter]
readRTSStats = do
  iscollected ← GhcStats.getRTSStatsEnabled
  if iscollected
    then ghcstats
    else return []
  where
    ghcstats :: IO [Counter]
    ghcstats = do
      -- need to run GC?
      rts ← GhcStats.getRTSStats
      let getrts = ghcval rts
      return [getrts (toInteger ∘ GhcStats.allocated_bytes, "bytesAllocated")
        ,getrts (toInteger ∘ GhcStats.max_live_bytes, "liveBytes")
        ,getrts (toInteger ∘ GhcStats.max_large_objects_bytes, "largeBytes")
        ,getrts (toInteger ∘ GhcStats.max_compact_bytes, "compactBytes")
        ,getrts (toInteger ∘ GhcStats.max_slop_bytes, "slopBytes")
        ,getrts (toInteger ∘ GhcStats.max_mem_in_use_bytes, "usedMemBytes")
        ,getrts (toInteger ∘ GhcStats.gc_cpu_ns, "gcCpuNs")
        ,getrts (toInteger ∘ GhcStats.gc_elapsed_ns, "gcElapsedNs")
        ,getrts (toInteger ∘ GhcStats.cpu_ns, "cpuNs")
        ,getrts (toInteger ∘ GhcStats.elapsed_ns, "elapsedNs")
        ,getrts (toInteger ∘ GhcStats.gcs, "gcNum")
        ,getrts (toInteger ∘ GhcStats.major_gcs, "gcMajorNum")
        ]
    ghcval :: GhcStats.RTSStats → ((GhcStats.RTSStats → Integer), Text) → Counter
    ghcval s (f, n) = Counter RTSStats n $ PureI (f s)
```

1.4.8 Cardano.BM.Counters.Dummy

This is a dummy definition of *readCounters* on platforms that do not support the 'proc' filesystem from which we would read the counters.

The only supported measurements are monotonic clock time and RTS statistics for now.

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters (TeeTrace _) = return []
readCounters (FilterTrace _) = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
readCounters (ObservableTrace tts) = foldrM (\(sel,fun) a →
  if any (≡ sel) tts
  then (fun >>= λxs → return $ a ++ xs)
  else return a) [] selectors
where
  selectors = [(MonotonicClock, getMonoClock)
    -- , (MemoryStats, readProcStatM)
    -- , (ProcessStats, readProcStats)
    -- , (IOStats, readProcIO)
    , (GhcRtsStats, readRTSStats)
  ]

```

1.4.9 Cardano.BM.Counters.Linux

we have to expand the *readMemStats* function to read full data from *proc*

```

readCounters :: SubTrace → IO [Counter]
readCounters NoTrace      = return []
readCounters Neutral      = return []
readCounters (TeeTrace _) = return []
readCounters (FilterTrace _) = return []
readCounters UntimedTrace = return []
readCounters DropOpening  = return []
readCounters (ObservableTrace tts) = foldrM (\(sel,fun) a →
  if any (≡ sel) tts
  then (fun >>= λxs → return $ a ++ xs)
  else return a) [] selectors
where
  selectors = [(MonotonicClock, getMonoClock)
    , (MemoryStats, readProcStatM)
    , (ProcessStats, readProcStats)
    , (IOStats, readProcIO)
  ]

```



```

pathProc :: FilePath
pathProc = "/proc/"
pathProcStat :: ProcessID → FilePath
pathProcStat pid = pathProc < / > (show pid) < / > "stat"
pathProcStatM :: ProcessID → FilePath
pathProcStatM pid = pathProc < / > (show pid) < / > "statm"
pathProcIO :: ProcessID → FilePath
pathProcIO pid = pathProc < / > (show pid) < / > "io"

```

Reading from a file in /proc/<pid >

```

readProcList :: FilePath → IO [Integer]
readProcList fp = do
  cs ← readFile fp
  return $ map (λs → maybe 0 id $ (readMaybe s :: Maybe Integer)) (words cs)

```

readProcStatM - /proc/<pid >/statm

```

/proc/[pid]/statm
  Provides information about memory usage, measured in pages. The columns are:
  size          (1) total program size
                  (same as VmSize in /proc/[pid]/status)
  resident      (2) resident set size
                  (same as VmRSS in /proc/[pid]/status)
  shared        (3) number of resident shared pages (i.e., backed by a file)
                  (same as RssFile+RssShmem in /proc/[pid]/status)
  text          (4) text (code)
  lib           (5) library (unused since Linux 2.6; always 0)
  data          (6) data + stack
  dt            (7) dirty pages (unused since Linux 2.6; always 0)

```

```

readProcStatM :: IO [Counter]
readProcStatM = do
  pid ← getProcessID
  ps0 ← readProcList (pathProcStatM pid)
  let ps = zip colnames ps0
      psUseful = filter (("unused" ≠) ∘ fst) ps
  return $ map (λ(n,i) → Counter MemoryCounter n (PureI i)) psUseful
where
  colnames :: [Text]
  colnames = ["size", "resident", "shared", "text", "unused", "data", "unused"]

```

readProcStats - //proc//<pid >//stat

```

/proc/[pid]/stat
  Status information about the process. This is used by ps(1). It is defined in the kernel source file
  fs/proc/array.c.

```

The fields, in order, with their proper scanf(3) format specifiers, are listed below. Whether or not certain of these fields display valid information is governed by a ptrace access mode `PTTRACE_MODE_READ_FSCREDS | PTTRACE_MODE_NOAUDIT` check (refer to `ptrace(2)`). If the check denies access,

then the field value is displayed as 0. The affected fields are indicated with the marking [PT].

- (1) pid %d
The process ID.
- (2) comm %s
The filename of the executable, in parentheses. This is visible whether or not the executable is swapped out.
- (3) state %c
One of the following characters, indicating process state:
 - R Running
 - S Sleeping in an interruptible wait
 - D Waiting in uninterruptible disk sleep
 - Z Zombie
 - T Stopped (on a signal) or (before Linux 2.6.33) trace stopped
 - t Tracing stop (Linux 2.6.33 onward)
 - W Paging (only before Linux 2.6.0)
 - X Dead (from Linux 2.6.0 onward)
 - x Dead (Linux 2.6.33 to 3.13 only)
 - K Wakekill (Linux 2.6.33 to 3.13 only)
 - W Waking (Linux 2.6.33 to 3.13 only)
 - P Parked (Linux 3.9 to 3.13 only)
- (4) ppid %d
The PID of the parent of this process.
- (5) pgrp %d
The process group ID of the process.
- (6) session %d
The session ID of the process.
- (7) tty_nr %d
The controlling terminal of the process. (The minor device number is contained in the combination of bits 31 to 20 and 7 to 0; the major device number is in bits 15 to 8.)
- (8) tpgid %d
The ID of the foreground process group of the controlling terminal of the process.
- (9) flags %u
The kernel flags word of the process. For bit meanings, see the PF_* defines in the Linux kernel source file include/linux/sched.h. Details depend on the kernel version.

The format for this field was %lu before Linux 2.6.
- (10) minflt %lu
The number of minor faults the process has made which have not required loading a memory page from disk.
- (11) cminflt %lu
The number of minor faults that the process's waited-for children have made.
- (12) majflt %lu
The number of major faults the process has made which have required loading a memory page from disk.

- (13) `cmajflt %lu`
The number of major faults that the process's waited-for children have made.
- (14) `utime %lu`
Amount of time that this process has been scheduled in user mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`). This includes guest time, `guest_time` (time spent running a virtual CPU, see below), so that applications that are not aware of the guest time field do not lose that time from their calculations.
- (15) `stime %lu`
Amount of time that this process has been scheduled in kernel mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (16) `cutime %ld`
Amount of time that this process's waited-for children have been scheduled in user mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`). (See also `times(2)`.) This includes guest time, `cguest_time` (time spent running a virtual CPU, see below).
- (17) `cstime %ld`
Amount of time that this process's waited-for children have been scheduled in kernel mode, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (18) `priority %ld`
(Explanation for Linux 2.6) For processes running a real-time scheduling policy (policy below; see `sched_setscheduler(2)`), this is the negated scheduling priority, minus one; that is, a number in the range -2 to -100, corresponding to real-time priorities 1 to 99. For processes running under a non-real-time scheduling policy, this is the raw nice value (`setpriority(2)`) as represented in the kernel. The kernel stores nice values as numbers in the range 0 (high) to 39 (low), corresponding to the user-visible nice range of -20 to 19.
- (19) `nice %ld`
The nice value (see `setpriority(2)`), a value in the range 19 (low priority) to -20 (high priority).
- (20) `num_threads %ld`
Number of threads in this process (since Linux 2.6). Before kernel 2.6, this field was hard coded to 0 as a placeholder for an earlier removed field.
- (21) `itrealvalue %ld`
The time in jiffies before the next `SIGALRM` is sent to the process due to an interval timer. Since kernel 2.6.17, this field is no longer maintained, and is hard coded as 0.
- (22) `starttime %llu`
The time the process started after system boot. In kernels before Linux 2.6, this value was expressed in jiffies. Since Linux 2.6, the value is expressed in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).

The format for this field was `%lu` before Linux 2.6.
- (23) `vsize %lu`
Virtual memory size in bytes.
- (24) `rss %ld`
Resident Set Size: number of pages the process has in real memory. This is just the pages which count toward text, data, or stack space. This does not include pages which have not been demand-loaded in, or which are swapped out.
- (25) `rsslim %lu`
Current soft limit in bytes on the `rss` of the process; see the description of `RLIMIT_RSS` in `getrlimit(2)`.
- (26) `startcode %lu [PT]`
The address above which program text can run.
- (27) `endcode %lu [PT]`
The address below which program text can run.

- (28) `startstack %lu [PT]`
The address of the start (i.e., bottom) of the stack.
- (29) `kstkesp %lu [PT]`
The current value of ESP (stack pointer), as found in the kernel stack page for the process.
- (30) `kstkeip %lu [PT]`
The current EIP (instruction pointer).
- (31) `signal %lu`
The bitmap of pending signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (32) `blocked %lu`
The bitmap of blocked signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (33) `sigignore %lu`
The bitmap of ignored signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (34) `sigcatch %lu`
The bitmap of caught signals, displayed as a decimal number. Obsolete, because it does not provide information on real-time signals; use `/proc/[pid]/status` instead.
- (35) `wchan %lu [PT]`
This is the "channel" in which the process is waiting. It is the address of a location in the kernel where the process is sleeping. The corresponding symbolic name can be found in `/proc/[pid]/wchan`.
- (36) `nswap %lu`
Number of pages swapped (not maintained).
- (37) `cnsnap %lu`
Cumulative nswap for child processes (not maintained).
- (38) `exit_signal %d (since Linux 2.1.22)`
Signal to be sent to parent when we die.
- (39) `processor %d (since Linux 2.2.8)`
CPU number last executed on.
- (40) `rt_priority %u (since Linux 2.5.19)`
Real-time scheduling priority, a number in the range 1 to 99 for processes scheduled under a real-time policy, or 0, for non-real-time processes (see `sched_setscheduler(2)`).
- (41) `policy %u (since Linux 2.5.19)`
Scheduling policy (see `sched_setscheduler(2)`). Decode using the `SCHED_*` constants in `linux/sched.h`.

The format for this field was `%lu` before Linux 2.6.22.
- (42) `delayacct_blkio_ticks %llu (since Linux 2.6.18)`
Aggregated block I/O delays, measured in clock ticks (centiseconds).
- (43) `guest_time %lu (since Linux 2.6.24)`
Guest time of the process (time spent running a virtual CPU for a guest operating system), measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (44) `cguest_time %ld (since Linux 2.6.24)`
Guest time of the process's children, measured in clock ticks (divide by `sysconf(_SC_CLK_TCK)`).
- (45) `start_data %lu (since Linux 3.3) [PT]`
Address above which program initialized and uninitialized (BSS) data are placed.
- (46) `end_data %lu (since Linux 3.3) [PT]`
Address below which program initialized and uninitialized (BSS) data are placed.

- (47) `start_brk %lu` (since Linux 3.3) [PT]
Address above which program heap can be expanded with `brk(2)`.
- (48) `arg_start %lu` (since Linux 3.5) [PT]
Address above which program command-line arguments (`argv`) are placed.
- (49) `arg_end %lu` (since Linux 3.5) [PT]
Address below program command-line arguments (`argv`) are placed.
- (50) `env_start %lu` (since Linux 3.5) [PT]
Address above which program environment is placed.
- (51) `env_end %lu` (since Linux 3.5) [PT]
Address below which program environment is placed.
- (52) `exit_code %d` (since Linux 3.5) [PT]
The thread's exit status in the form reported by `waitpid(2)`.

```
readProcStats :: IO [Counter]
```

```
readProcStats = do
```

```
    pid ← getProcessID
```

```
    ps0 ← readProcList (pathProcStat pid)
```

```
    let ps = zip colnames ps0
```

```
        psUseful = filter (("unused"  $\neq$ )  $\circ$  fst) ps
```

```
    return $ map ( $\lambda(n,i) \rightarrow$  Counter StatInfo  $n$  (PureI  $i$ )) psUseful
```

```
where
```

```
colnames :: [Text]
```

```
colnames = [ "pid", "unused", "unused", "ppid", "pgrp", "session", "ttynr", "tpgid", "flags", "minflt",
"cminflt", "majflt", "cmajflt", "utime", "stime", "cutime", "cstime", "priority", "nice", "num",
"itrealvalue", "starttime", "vsize", "rss", "rsslim", "startcode", "endcode", "startstack",
"signal", "blocked", "sigignore", "sigcatch", "wchan", "nswap", "cnsnap", "exitsignal", "proc",
"policy", "blkio", "guesttime", "cguesttime", "startdata", "enddata", "startbrk", "argstart",
"envend", "exitcode"
]
```

readProcIO - `//proc//<pid >//io`

```
/proc/[pid]/io (since kernel 2.6.20)
```

This file contains I/O statistics for the process, for example:

```
# cat /proc/3828/io
rchar: 323934931
wchar: 323929600
syscr: 632687
syscw: 632675
read_bytes: 0
write_bytes: 323932160
cancelled_write_bytes: 0
```

The fields are as follows:

`rchar`: characters read

The number of bytes which this task has caused to be read from storage. This is simply the sum of bytes which this process passed to `read(2)` and similar system calls. It includes things such as terminal I/O and is unaffected by whether or not actual physical disk I/O was required (the read might have been satisfied from pagecache).

wchar: characters written
The number of bytes which this task has caused, or shall cause to be written to disk. Similar caveats apply here as with rchar.

syscr: read syscalls
Attempt to count the number of read I/O operations—that is, system calls such as read(2) and pread(2).

syscw: write syscalls
Attempt to count the number of write I/O operations—that is, system calls such as write(2) and pwrite(2).

read_bytes: bytes read
Attempt to count the number of bytes which this process really did cause to be fetched from the storage layer. This is accurate for block-backed filesystems.

write_bytes: bytes written
Attempt to count the number of bytes which this process caused to be sent to the storage layer.

cancelled_write_bytes:
The big inaccuracy here is truncate. If a process writes 1MB to a file and then deletes the file, it will in fact perform no writeout. But it will have been accounted as having caused 1MB of write. In other words: this field represents the number of bytes which this process caused to not happen, by truncating pagecache. A task can cause "negative" I/O too. If this task truncates some dirty pagecache, some I/O which another task has been accounted for (in its write_bytes) will not be happening.

Note: In the current implementation, things are a bit racy on 32-bit systems: if process A reads process B's /proc/[pid]/io while process B is updating one of these 64-bit counters, process A could see an intermediate result.

Permission to access this file is governed by a ptrace access mode PTRACE_MODE_READ_FSCREDS check; see ptrace(2).

```
readProcIO :: IO [Counter]
readProcIO = do
  pid ← getProcessID
  ps0 ← readProcList (pathProcIO pid)
  let ps = zip3 colnames ps0 units
  return $ map (\(n,i,u) → Counter IOCounter n (u i)) ps
where
  colnames :: [Text]
  colnames = [ "rchar", "wchar", "syscr", "syscw", "rbytes", "wbytes", "cxwbytes" ]
  units = [ Bytes, Bytes, PureI, PureI, Bytes, Bytes, Bytes ]
```

1.4.10 Cardano.BM.Data.Aggregated

Measurable

A **Measurable** may consist of different types of values.

```
data Measurable = Microseconds Integer
                | Seconds Integer
                | Bytes Integer
                | PureI Integer
                | PureD Double
deriving (Eq, Ord, Generic, ToJSON)
```

Measurable can be transformed to an integral value.

```
getInteger :: Measurable → Integer
getInteger (Microseconds a) = a
getInteger (Seconds a)      = a
getInteger (Bytes a)        = a
getInteger (PureI a)        = a
getInteger (PureD a)        = round a
```

Measurable can be transformed to a rational value.

```
getDouble :: Measurable → Double
getDouble (Microseconds a) = fromInteger a
getDouble (Seconds a)     = fromInteger a
getDouble (Bytes a)       = fromInteger a
getDouble (PureI a)       = fromInteger a
getDouble (PureD a)       = a
```

It is a numerical value, thus supports functions to operate on numbers.

```
instance Num Measurable where
  (+) (Microseconds a) (Microseconds b) = Microseconds (a + b)
  (+) (Seconds a) (Seconds b) = Seconds (a + b)
  (+) (Bytes a) (Bytes b) = Bytes (a + b)
  (+) (PureI a) (PureI b) = PureI (a + b)
  (+) (PureD a) (PureD b) = PureD (a + b)
  (+) _ _ = error "Trying to add values with different units"
  (*) (Microseconds a) (Microseconds b) = Microseconds (a * b)
  (*) (Seconds a) (Seconds b) = Seconds (a * b)
  (*) (Bytes a) (Bytes b) = Bytes (a * b)
  (*) (PureI a) (PureI b) = PureI (a * b)
  (*) (PureD a) (PureD b) = PureD (a * b)
  (*) _ _ = error "Trying to multiply values with different units"
  abs (Microseconds a) = Microseconds (abs a)
  abs (Seconds a) = Seconds (abs a)
  abs (Bytes a) = Bytes (abs a)
  abs (PureI a) = PureI (abs a)
  abs (PureD a) = PureD (abs a)
  signum (Microseconds a) = Microseconds (signum a)
  signum (Seconds a) = Seconds (signum a)
  signum (Bytes a) = Bytes (signum a)
  signum (PureI a) = PureI (signum a)
  signum (PureD a) = PureD (signum a)
  negate (Microseconds a) = Microseconds (negate a)
  negate (Seconds a) = Seconds (negate a)
  negate (Bytes a) = Bytes (negate a)
  negate (PureI a) = PureI (negate a)
  negate (PureD a) = PureD (negate a)
```

```
fromInteger = PureI
```

Pretty printing of **Measurable**.

```
instance Show Measurable where
  show = showSI

showUnits :: Measurable → String
showUnits (Microseconds _) = " s"
showUnits (Seconds _) = " s"
showUnits (Bytes _) = " B"
showUnits (PureI _) = ""
showUnits (PureD _) = ""

-- show in S.I. units
showSI :: Measurable → String
showSI (Microseconds a) = show (fromFloatDigits ((fromInteger a) / (1000000 :: Float))) ++
  showUnits (Seconds a)
showSI v@(Seconds a) = show a ++ showUnits v
showSI v@(Bytes a) = show a ++ showUnits v
showSI v@(PureI a) = show a ++ showUnits v
showSI v@(PureD a) = show a ++ showUnits v
```

Stats

```
data Stats = Stats {
  flast :: Measurable,
  fmin :: Measurable,
  fmax :: Measurable,
  fcount :: Integer,
  fsum_A :: Double,
  fsum_B :: Double
} deriving (Eq, Generic, ToJSON, Show)

meanOfStats :: Stats → Double
meanOfStats s = fsum_A s

stdevOfStats :: Stats → Double
stdevOfStats s =
  if fcount s < 2
  then 0
  else sqrt $ (fsum_B s) / (fromInteger $ (fcount s) - 1)
```

instance Semigroup Stats disabled for the moment, because not needed.

We use a parallel algorithm to update the estimation of mean and variance from two sample statistics. (see https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Parallel_algorithm)


```

instance Semigroup Stats where
  (<>) a b = let counta = fcount a
             countb = fcount b
             newcount = counta + countb
             delta = fsum_A b - fsum_A a
             in
    Stats {flast = flast b -- right associative
           ,fmin   = min (fmin a) (fmin b)
           ,fmax   = max (fmax a) (fmax b)
           ,fcount = newcount
           ,fsum_A = fsum_A a + (delta / fromInteger newcount)
           ,fsum_B = fsum_B a + fsum_B b + (delta * delta) * (fromInteger (counta * countb) / fromInteger newcount)
           }

```

```

stats2Text :: Stats → Text
stats2Text s@(Stats slast smin smax scount _) =
  pack $
    "{ last = " ++ show slast ++
    ", min = " ++ show smin ++
    ", max = " ++ show smax ++
    ", mean = " ++ show (meanOfStats s) ++ showUnits slast ++
    ", std-dev = " ++ show (stdevOfStats s) ++
    ", count = " ++ show scount ++
    " }"

```

Exponentially Weighted Moving Average (EWMA)

```

data EWMA = EmptyEWMA {alpha :: Double}
  | EWMA {alpha :: Double
         ,avg :: Measurable
         } deriving (Show, Eq, Generic, ToJSON)

```

Aggregated

```

data Aggregated = AggregatedStats Stats
  | AggregatedEWMA EWMA
deriving (Eq, Generic, ToJSON)

```

instance Semigroup **Aggregated** disabled for the moment, because not needed.

```

instance Semigroup Aggregated where
  (<>) (AggregatedStats a) (AggregatedStats b) =
    AggregatedStats (a <> b)
  (<>) _ _ = error "Cannot combine different objects"

```

```

singletonStats :: Measurable → Aggregated
singletonStats a =
  let stats = Stats {flast = a
                    ,fmin      = a
                    ,fmax      = a
                    ,fcount = 1
                    ,fsum_A = getDouble a
                    ,fsum_B = 0
                    }
  in
    AggregatedStats stats

```

```

instance Show Aggregated where
  show (AggregatedStats astats) =
    "{ stats = " ++ show astats ++ " }"
  show (AggregatedEWMA a) = show a

```

1.4.11 Cardano.BM.Data.Backend

Accepts a NamedLogItem

Instances of this type class accept a **NamedLogItem** and deal with it.

```

class IsEffectuator t where
  effectuate :: t → NamedLogItem → IO ()
  effectuatefrom :: forall s ◦ (IsEffectuator s) ⇒ t → NamedLogItem → s → IO ()
  default effectuatefrom :: forall s ◦ (IsEffectuator s) ⇒ t → NamedLogItem → s → IO ()
  effectuatefrom t nli _ = effectuate t nli

```

Declaration of a Backend

A backend is life-cycle managed, thus can be *realized* and *unrealized*.

```

class (IsEffectuator t) ⇒ IsBackend t where
  typeof    :: t → BackendKind
  realize    :: Configuration → IO t
  realizefrom :: forall s ◦ (IsEffectuator s) ⇒ Trace IO → s → IO t
  default realizefrom :: forall s ◦ (IsEffectuator s) ⇒ Trace IO → s → IO t
  realizefrom (ctx, _) _ = realize (configuration ctx)
  unrealize :: t → IO ()

```

Backend

This data structure for a backend defines its behaviour as an **IsEffectuator** when processing an incoming message, and as an **IsBackend** for unrealizing the backend.

```

data Backend = MkBackend
  { bEffectuate :: NamedLogItem → IO ()

```

```
,bUnrealize :: IO ()
}
```

1.4.12 Cardano.BM.Data.Configuration

Data structure to help parsing configuration files.

Representation

```
type Port = Int
data Representation = Representation
  { minSeverity    :: Severity
  , rotation       :: RotationParameters
  , setupScribes   :: [ScribeDefinition]
  , defaultScribes :: [(ScribeKind, Text)]
  , setupBackends  :: [BackendKind]
  , defaultBackends :: [BackendKind]
  , hasEKG         :: Maybe Port
  , hasGUI         :: Maybe Port
  , options        :: HM.HashMap Text Object
  }
deriving (Generic, Show, ToJSON, FromJSON)
```

parseRepresentation

```
parseRepresentation :: FilePath → IO Representation
parseRepresentation fp = do
  repr :: Representation ← decodeFileThrow fp
  return $ implicit_fill_representation repr
```

after parsing the configuration representation we implicitly correct it.

```
implicit_fill_representation :: Representation → Representation
implicit_fill_representation =
  remove_ekgview_if_not_defined ∘
  filter_duplicates_from_backends ∘
  filter_duplicates_from_scribes ∘
  union_setup_and_usage_backends ∘
  add_ekgview_if_port_defined ∘
  add_katip_if_any_scribes
where
  filter_duplicates_from_backends r =
    r {setupBackends = mkUniq $ setupBackends r}
  filter_duplicates_from_scribes r =
    r {setupScribes = mkUniq $ setupScribes r}
  union_setup_and_usage_backends r =
```

```

    r {setupBackends = setupBackends r <> defaultBackends r}
remove_ekgview_if_not_defined r =
  case hasEKG r of
    Nothing → r {defaultBackends = filter (λbk → bk ≠ EKGViewBK) (defaultBackends r)
      ,setupBackends = filter (λbk → bk ≠ EKGViewBK) (setupBackends r)
    }
    Just _ → r
add_ekgview_if_port_defined r =
  case hasEKG r of
    Nothing → r
    Just _ → r {setupBackends = setupBackends r <> [EKGViewBK]}
add_katip_if_any_scribes r =
  if (any ¬ [null $ setupScribes r, null $ defaultScribes r])
  then r {setupBackends = setupBackends r <> [KatipBK]}
  else r
mkUniq :: Ord a ⇒ [a] → [a]
mkUniq = Set.toList ∘ Set.fromList

```

1.4.13 Cardano.BM.Data.Counter

Counter

```

data Counter = Counter
  { cType :: CounterType
  , cName :: Text
  , cValue :: Measurable
  }
  deriving (Eq, Show, Generic, ToJSON)

data CounterType = MonotonicClockTime
  | MemoryCounter
  | StatInfo
  | IOCounter
  | CpuCounter
  | RTStats
  deriving (Eq, Show, Generic, ToJSON)

instance ToJSON Microsecond where
  toJSON = toJSON ∘ toMicroseconds
  toEncoding = toEncoding ∘ toMicroseconds

```

Names of counters

```

nameCounter :: Counter → Text
nameCounter (Counter MonotonicClockTime _ _) = "Time-interval"
nameCounter (Counter MemoryCounter _ _) = "Mem"
nameCounter (Counter StatInfo _ _) = "Stat"

```

```

nameCounter (Counter IOCounter    _ _) = "IO"
nameCounter (Counter CpuCounter   _ _) = "Cpu"
nameCounter (Counter RTStats     _ _) = "RTS"

```

CounterState

```

data CounterState = CounterState {
  csIdentifier :: Unique
  ,csCounters :: [Counter]
}
deriving (Generic, ToJSON)
instance ToJSON Unique where
  toJSON = toJSON ∘ hashUnique
  toEncoding = toEncoding ∘ hashUnique
instance Show CounterState where
  show cs = (show ∘ hashUnique) (csIdentifier cs)
  <> " => " <> (show $ csCounters cs)

```

Difference between counters

```

diffCounters :: [Counter] → [Counter] → [Counter]
diffCounters openings closings =
  getCountersDiff openings closings
where
  getCountersDiff :: [Counter]
    → [Counter]
    → [Counter]
  getCountersDiff as bs =
    let
      getName counter = nameCounter counter <> cName counter
      asNames = map getName as
      aPairs = zip asNames as
      bsNames = map getName bs
      bs' = zip bsNames bs
      bPairs = HM.fromList bs'
    in
      catMaybes $ (flip map) aPairs $ λ(name, Counter _ _ startValue) →
        case HM.lookup name bPairs of
          Nothing    → Nothing
          Just counter → let endValue = cValue counter
                        in Just counter {cValue = endValue − startValue}

```

1.4.14 Cardano.BM.Data.LogItem

LoggerName

```
type LoggerName = Text
```

NamedLogItem

```
type NamedLogItem = LogNamed LogObject
```

LogItem

```
TODO liPayload :: ToObject
```

```
data LogItem = LogItem
  { liSelection :: LogSelection
  , liSeverity :: Severity
  , liPayload :: Text -- TODO should become ToObject
  } deriving (Show, Generic, ToJSON)
```

```
data LogSelection =
  Public -- only to public logs.
  | PublicUnsafe -- only to public logs, not console.
  | Private -- only to private logs.
  | Both -- to public and private logs.
  deriving (Show, Generic, ToJSON, FromJSON)
```

LogObject

```
data LogPrims = LogMessage LogItem
  | LogValue Text Measurable
  deriving (Generic, Show, ToJSON)
data LogObject = LP LogPrims
  | ObserveOpen CounterState
  | ObserveDiff CounterState
  | ObserveClose CounterState
  | AggregatedMessage [(Text, Aggregated)]
  | KillPill
  deriving (Generic, Show, ToJSON)
```

LogNamed

A **LogNamed** contains of a context name and some log item.

```
data LogNamed item = LogNamed
  { lnName :: LoggerName
  , lnItem :: item
  } deriving (Show)
deriving instance Generic item => Generic (LogNamed item)
deriving instance (ToJSON item, Generic item) => ToJSON (LogNamed item)
```

1.4.15 Cardano.BM.Data.Observable**ObservableInstance**

```
data ObservableInstance = MonotonicClock
  | MemoryStats
  | ProcessStats
  | IOStats
  | GhcRtsStats
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON, Read)
```

1.4.16 Cardano.BM.Data.Output**OutputKind**

```
data OutputKind = TVarList (STM.TVar [LogObject])
  | TVarListNamed (STM.TVar [LogNamed LogObject])
  deriving (Eq)
```

ScribeKind

This identifies katip's scribes by type.

```
data ScribeKind = FileTextSK
  | FileJsonSK
  | StdoutSK
  | StderrSK
  deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

ScribeId

A scribe is identified by **ScribeKind** x *Filename*

```
type ScribeId = Text-- (ScribeKind :: Filename)
```

ScribeDefinition

This identifies katip's scribes by type.

```
data ScribeDefinition = ScribeDefinition
  { scKind :: ScribeKind
  , scName :: Text
  , scRotation :: Maybe RotationParameters
  }
deriving (Generic, Eq, Ord, Show, FromJSON, ToJSON)
```

1.4.17 Cardano.BM.Data.Severity

Severity

The intended meaning of severity codes:

Debug *detailed information about values and decision flow* **Info** general information of events; progressing properly *Notice needs attention; something \rightarrow progressing properly* **Warning** may continue into an error condition if continued *Error unexpected set of event or condition occurred* **Critical** error condition causing degrade of operation *Alert a subsystem is no longer operating correctly, likely requires manual* at this point, the system can never progress without additional intervention

We were informed by the Syslog taxonomy: https://en.wikipedia.org/wiki/Syslog#Severity_level

```
data Severity = Debug
  | Info
  | Notice
  | Warning
  | Error
  | Critical
  | Alert
  | Emergency
  deriving (Show, Eq, Ord, Generic, ToJSON, Read)

instance FromJSON Severity where
  parseJSON = withText "severity" $ \case
    "Debug"    → pure Debug
    "Info"     → pure Info
    "Notice"   → pure Notice
    "Warning"  → pure Warning
    "Error"    → pure Error
    "Critical" → pure Critical
    "Alert"    → pure Alert
    "Emergency" → pure Emergency
    _          → pure Info -- catch all
```

1.4.18 Cardano.BM.Data.SubTrace

SubTrace

```
data NameSelector = Exact Text | StartsWith Text | EndsWith Text | Contains Text
  | Named Text -- LP (LogValue (name value))
```



```

    deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
data NameOperator = Drop NameSelector | Unhide NameSelector
    deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)
data SubTrace = Neutral
    | UntimedTrace
    | NoTrace
    | TeeTrace LoggerName
    | FilterTrace [NameOperator]
    | DropOpening
    | ObservableTrace [ObservableInstance]
    deriving (Generic, Show, FromJSON, ToJSON, Read, Eq)

```

1.4.19 Cardano.BM.Data.Trace

Trace

A **Trace** consists of a **TraceContext** and a **TraceNamed** in *m*.

```
type Trace m = (TraceContext, TraceNamed m)
```

TraceNamed

A **TraceNamed** is a specialized **Contravariant** of type **NamedLogItem**, a **LogNamed** with payload **LogObject**.

```
type TraceNamed m = BaseTrace m (NamedLogItem)
```

TraceContext

We keep the context's name and a reference to the **Configuration** in the **TraceContext**.

```

data TraceContext = TraceContext
    { loggerName :: LoggerName
    , configuration :: Configuration
    , tracetype   :: SubTrace
    , minSeverity :: Severity
    , shutdown    :: IO ()
    }

```

1.4.20 Cardano.BM.Configuration

see **Cardano.BM.Configuration.Model** for the implementation.

```

getOptionOrDefault :: CM.Configuration → Text → Text → IO (Text)
getOptionOrDefault cg name def = do
    opt ← CM.getOption cg name
    case opt of
        Nothing → return def
        Just o  → return o

```

1.4.21 Cardano.BM.Configuration.Model

Configuration.Model

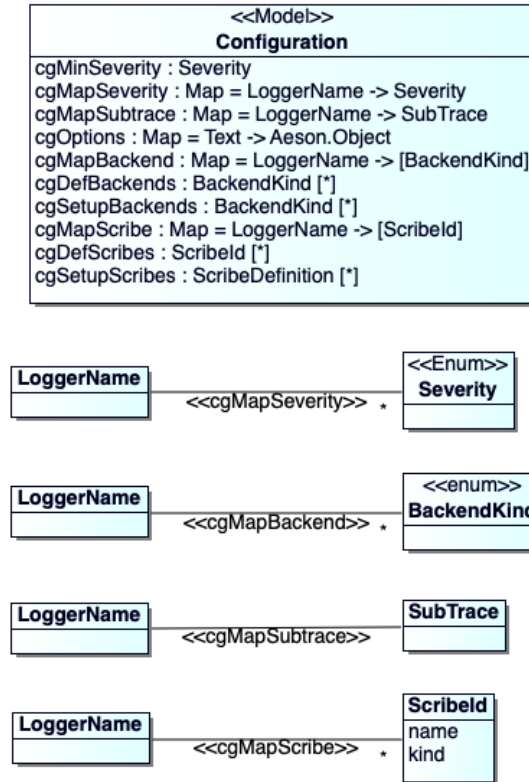


Figure 1.4: Configuration model

```

type ConfigurationMVar = MVar ConfigurationInternal
newtype Configuration = Configuration
    {getCG :: ConfigurationMVar}
-- Our internal state; see -"Configuration model"-
data ConfigurationInternal = ConfigurationInternal
    {cgMinSeverity    :: Severity
    ,cgMapSeverity    :: HM.HashMap LoggerName Severity
    ,cgMapSubtrace    :: HM.HashMap LoggerName SubTrace
    ,cgOptions        :: HM.HashMap Text Object
    ,cgMapBackend     :: HM.HashMap LoggerName [BackendKind]
    ,cgDefBackendKs   :: [BackendKind]
  
```

```

-- backends that will be used if a set of backends for the
-- specific loggename is not set
,cgSetupBackends :: [BackendKind]
-- backends to setup; every backend to be used must have
-- been declared here
,cgMapScribe      :: HM.HashMap LoggerName [ScribeId]
-- katip scribes that will be used for the specific loggename
,cgDefScribes     :: [ScribeId]
-- katip scribes that will be used if a set of scribes for the
-- specific loggename is not set
,cgSetupScribes  :: [ScribeDefinition]
-- katip scribes to setup; every scribe to be used must have
-- been declared here
,cgMapAggregatedKind :: HM.HashMap LoggerName AggregatedKind
-- kind of Aggregated that will be used for the specific loggename
,cgDefAggregatedKind :: AggregatedKind
-- kind of Aggregated that will be used if a set of scribes for the
-- specific loggename is not set
,cgPortEKG       :: Int
-- port for EKG server
,cgPortGUI       :: Int
-- port for changes at runtime (NOT IMPLEMENTED YET)
} deriving (Show, Eq)

```

Backends configured in the Switchboard

For a given context name return the list of backends configured, or, in case no such configuration exists, return the default backends.

```

getBackends :: Configuration → LoggerName → IO [BackendKind]
getBackends configuration name =
  withMVar (getCG configuration) $ \cg → do
    let outs = HM.lookup name (cgMapBackend cg)
    case outs of
      Nothing → do
        return (cgDefBackendKs cg)
      Just os → return os

getDefaultBackends :: Configuration → IO [BackendKind]
getDefaultBackends configuration =
  withMVar (getCG configuration) $ \cg → do
    return (cgDefBackendKs cg)

setDefaultBackends :: Configuration → [BackendKind] → IO ()
setDefaultBackends configuration bes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefBackendKs = bes}

setBackends :: Configuration → LoggerName → Maybe [BackendKind] → IO ()
setBackends configuration name be = do

```

```
cg ← takeMVar (getCG configuration)
putMVar (getCG configuration) $ cg {cgMapBackend = HM.alter (\_ → be) name (cgMapBackend cg)}
```

Backends to be setup by the Switchboard

Defines the list of **Backends** that need to be setup by the **Switchboard**.

```
setSetupBackends :: Configuration → [BackendKind] → IO ()
setSetupBackends configuration bes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgSetupBackends = bes}
getSetupBackends :: Configuration → IO [BackendKind]
getSetupBackends configuration =
  withMVar (getCG configuration) $ λcg →
    return $ cgSetupBackends cg
```

Scribes configured in the Log backend

For a given context name return the list of scribes to output to, or, in case no such configuration exists, return the default scribes to use.

```
getScribes :: Configuration → LoggerName → IO [ScribeId]
getScribes configuration name =
  withMVar (getCG configuration) $ λcg → do
    let outs = HM.lookup name (cgMapScribe cg)
    case outs of
      Nothing → do
        return (cgDefScribes cg)
      Just os → return $ os
setScribes :: Configuration → LoggerName → Maybe [ScribeId] → IO ()
setScribes configuration name scribes = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapScribe = HM.alter (\_ → scribes) name (cgMapScribe cg)}
setDefaultScribes :: Configuration → [ScribeId] → IO ()
setDefaultScribes configuration scs = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefScribes = scs}
```

Scribes to be setup in the Log backend

Defines the list of *Scribes* that need to be setup in the **Log** backend.

```
setSetupScribes :: Configuration → [ScribeDefinition] → IO ()
setSetupScribes configuration sds = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgSetupScribes = sds}
```

```

getSetupScribes :: Configuration → IO [ScribeDefinition]
getSetupScribes configuration =
  withMVar (getCG configuration) $ λcg → do
    return $ cgSetupScribes cg

```

AggregatedKind to define the type of measurement

For a given context name return its *AggregatedKind* or in case no such configuration exists, return the default *AggregatedKind* to use.

```

getAggregatedKind :: Configuration → LoggerName → IO AggregatedKind
getAggregatedKind configuration name =
  withMVar (getCG configuration) $ λcg → do
    let outs = HM.lookup name (cgMapAggregatedKind cg)
    case outs of
      Nothing → do
        return (cgDefAggregatedKind cg)
      Just os → return $ os

setDefaultAggregatedKind :: Configuration → AggregatedKind → IO ()
setDefaultAggregatedKind configuration defAK = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgDefAggregatedKind = defAK}

setAggregatedKind :: Configuration → LoggerName → Maybe AggregatedKind → IO ()
setAggregatedKind configuration name ak = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapAggregatedKind = HM.alter (\_ → ak) name (cgMapAggregatedKind cg)}

```

Access port numbers of EKG, GUI

```

getEKGport :: Configuration → IO Int
getEKGport configuration =
  withMVar (getCG configuration) $ λcg → do
    return $ cgPortEKG cg

setEKGport :: Configuration → Int → IO ()
setEKGport configuration port = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgPortEKG = port}

getGUIport :: Configuration → IO Int
getGUIport configuration =
  withMVar (getCG configuration) $ λcg → do
    return $ cgPortGUI cg

setGUIport :: Configuration → Int → IO ()
setGUIport configuration port = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgPortGUI = port}

```

Options

```
getOption :: Configuration → Text → IO (Maybe Text)
getOption configuration name = do
  withMVar (getCG configuration) $ \cg →
    case HM.lookup name (cgOptions cg) of
      Nothing → return Nothing
      Just o → return $ Just $ pack $ show o
```

Global setting of minimum severity

```
minSeverity :: Configuration → IO Severity
minSeverity configuration = withMVar (getCG configuration) $ \cg →
  return $ cgMinSeverity cg
setMinSeverity :: Configuration → Severity → IO ()
setMinSeverity configuration sev = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMinSeverity = sev}
```

Relation of context name to minimum severity

```
inspectSeverity :: Configuration → Text → IO (Maybe Severity)
inspectSeverity configuration name = do
  withMVar (getCG configuration) $ \cg →
    return $ HM.lookup name (cgMapSeverity cg)
setSeverity :: Configuration → Text → Maybe Severity → IO ()
setSeverity configuration name sev = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapSeverity = HM.alter (\_ → sev) name (cgMapSeverity cg)}
```

Relation of context name to SubTrace

A new context may contain a different type of **Trace**. The function **appendName** (Enter new named context) will look up the **SubTrace** for the context's name.

```
findSubTrace :: Configuration → Text → IO (Maybe SubTrace)
findSubTrace configuration name = do
  withMVar (getCG configuration) $ \cg →
    return $ HM.lookup name (cgMapSubtrace cg)
setSubTrace :: Configuration → Text → Maybe SubTrace → IO ()
setSubTrace configuration name trafo = do
  cg ← takeMVar (getCG configuration)
  putMVar (getCG configuration) $ cg {cgMapSubtrace = HM.alter (\_ → trafo) name (cgMapSubtrace cg)}
```

Parse configuration from file

Parse the configuration into an internal representation first. Then, fill in **Configuration** after refinement.

```

setup :: FilePath → IO Configuration
setup fp = do
  r ← R.parseRepresentation fp
  cgreg ← newEmptyMVar
  let mapseverity = HM.lookup "mapSeverity" (R.options r)
      mapbackends = HM.lookup "mapBackends" (R.options r)
      mapsubtrace = HM.lookup "mapSubtrace" (R.options r)
      mapscribes = HM.lookup "mapScribes" (R.options r)
      mapAggregatedKinds = HM.lookup "mapAggregatedkinds" (R.options r)
  putMVar cgreg $ ConfigurationInternal
    { cgMinSeverity = R.minSeverity r
    , cgMapSeverity = parseSeverityMap mapseverity
    , cgMapSubtrace = parseSubtraceMap mapsubtrace
    , cgOptions = R.options r
    , cgMapBackend = parseBackendMap mapbackends
    , cgDefBackendKs = R.defaultBackends r
    , cgSetupBackends = R.setupBackends r
    , cgMapScribe = parseScribeMap mapscribes
    , cgDefScribes = r_defaultScribes r
    , cgSetupScribes = R.setupScribes r
    , cgMapAggregatedKind = parseAggregatedKindMap mapAggregatedKinds
    , cgDefAggregatedKind = StatsAK
    , cgPortEKG = r_hasEKG r
    , cgPortGUI = r_hasGUI r
    }
  return $ Configuration cgreg
where
  parseSeverityMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text Severity
  parseSeverityMap Nothing = HM.empty
  parseSeverityMap (Just hmv) = HM.mapMaybe mkSeverity hmv
  mkSeverity (String s) = Just (read (unpack s) :: Severity)
  mkSeverity _ = Nothing

  parseBackendMap Nothing = HM.empty
  parseBackendMap (Just hmv) = HM.map mkBackends hmv
  mkBackends (Array bes) = catMaybes $ map mkBackend $ Vector.toList bes
  mkBackends _ = []
  mkBackend (String s) = Just (read (unpack s) :: BackendKind)
  mkBackend _ = Nothing

  parseScribeMap Nothing = HM.empty
  parseScribeMap (Just hmv) = HM.map mkScribes hmv
  mkScribes (Array scs) = catMaybes $ map mkScribe $ Vector.toList scs
  mkScribes (String s) = [(s :: ScribeId)]
  mkScribes _ = []

```

```

mkScribe (String s) = Just (s :: ScribeId)
mkScribe _ = Nothing

parseSubtraceMap :: Maybe (HM.HashMap Text Value) → HM.HashMap Text SubTrace
parseSubtraceMap Nothing = HM.empty
parseSubtraceMap (Just hmv) = HM.mapMaybe mkSubtrace hmv
mkSubtrace (String s) = Just (read (unpack s) :: SubTrace)
mkSubtrace (Object hm) = mkSubtrace' (HM.lookup "tag" hm) (HM.lookup "contents" hm)
mkSubtrace _ = Nothing
mkSubtrace' Nothing _ = Nothing
mkSubtrace' _ Nothing = Nothing
mkSubtrace' (Just (String tag)) (Just (Array cs)) =
  if tag == "ObservableTrace"
  then Just $ ObservableTrace $ map (λ(String s) → (read (unpack s) :: ObservableInstance)) $ Vector.toList cs
  else Nothing
mkSubtrace' _ _ = Nothing

r_hasEKG r = case (R.hasEKG r) of
  Nothing → 0
  Just p → p
r_hasGUI r = case (R.hasGUI r) of
  Nothing → 0
  Just p → p
r_defaultScribes r = map (λ(k,n) → pack (show k) <> " :: " <> n) (R.defaultScribes r)
parseAggregatedKindMap Nothing = HM.empty
parseAggregatedKindMap (Just hmv) =
  let
    listv = HM.toList hmv
    mapAggregatedKind = HM.fromList $ catMaybes $ map mkAggregatedKind listv
  in
    mapAggregatedKind
mkAggregatedKind (name, String s) = Just (name, read (unpack s) :: AggregatedKind)
mkAggregatedKind _ = Nothing

```

Setup empty configuration

```

empty :: IO Configuration
empty = do
  cgreg ← newEmptyMVar
  putMVar cgreg $ ConfigurationInternal Debug HM.empty HM.empty HM.empty HM.empty [ ] [ ] HM.empty [ ]
  return $ Configuration cgreg

```

1.4.22 Cardano.BM.Output.Switchboard

Switchboard

```

type SwitchboardMVar = MVar SwitchboardInternal
newtype Switchboard = Switchboard

```



```

    {getSB :: SwitchboardMVar}
data SwitchboardInternal = SwitchboardInternal
    {sbQueue :: TBQ.TBQueue NamedLogItem
    ,sbDispatch :: Async.Async ()
    }

```

Trace that forwards to the **Switchboard**

Every **Trace** ends in the **Switchboard** which then takes care of dispatching the messages to outputs

```

mainTrace :: Switchboard → TraceNamed IO
mainTrace sb = BaseTrace.BaseTrace $ Op $ λlognamed → do
    effectuate sb lognamed

```

Process incoming messages

Incoming messages are put into the queue, and then processed by the dispatcher.
The queue is initialized and the message dispatcher launched.

```

instance IsEffectuator Switchboard where
    effectuate switchboard item = do
        let writequeue :: TBQ.TBQueue NamedLogItem → NamedLogItem → IO ()
        writequeue q i = do
            nocapacity ← atomically $ TBQ.isFullTBQueue q
            if nocapacity
            then return ()
            else atomically $ TBQ.writeTBQueue q i
        withMVar (getSB switchboard) $ λsb →
            writequeue (sbQueue sb) item

```

Switchboard implements **Backend** functions

Switchboard is an **Declaration of a Backend**

```

instance IsBackend Switchboard where
    typeof _ = SwitchboardBK
    realize cfg =
        let spawnDispatcher :: Configuration → [(BackendKind, Backend)] → TBQ.TBQueue NamedLogItem → IO ()
        spawnDispatcher config backends queue =
            let sendMessage nli befilter = do
                selectedBackends ← getBackends config (lnName nli)
                let selBEs = befilter selectedBackends
                forM_ backends $ λ(bek, be) →
                    when (bek ∈ selBEs) (bEffectuate be $ nli)
            qProc = do
                nli ← atomically $ TBQ.readTBQueue queue

```

```

    case lnItem nli of
      KillPill →
        forM_ backends (λ(_, be) → bUnrealize be)
      AggregatedMessage _ → do
        sendMessage nli (filter (≠ AggregationBK))
        qProc
        _ → sendMessage nli id >> qProc
  in
    Async.async qProc
in do
  q ← atomically $ TBQ.newTBQueue 2048
  sbref ← newEmptyMVar
  putMVar sbref $ SwitchboardInternal q $ error "uninitialized dispatcher"
  let sb :: Switchboard = Switchboard sbref
  backends ← getSetupBackends cfg
  bs ← setupBackends backends cfg sb []
  dispatcher ← spawnDispatcher cfg bs q
  modifyMVar sbref $ λsbInternal → return $ sbInternal {sbDispatch = dispatcher}
  return sb
unrealize switchboard = do
  let clearMVar :: MVar a → IO ()
      clearMVar = void ∘ tryTakeMVar
  (dispatcher, queue) ← withMVar (getSB switchboard) (λsb → return (sbDispatch sb, sbQueue sb))
  -- send terminating item to the queue
  atomically $ TBQ.writeTBQueue queue $ LogNamed "kill.switchboard" KillPill
  -- wait for the dispatcher to exit
  res ← Async.waitCatch dispatcher
  either throwM return res
  (clearMVar ∘ getSB) switchboard

```

Realizing the backends according to configuration

```

setupBackends :: [BackendKind]
  → Configuration
  → Switchboard
  → [(BackendKind, Backend)]
  → IO [(BackendKind, Backend)]
setupBackends [] _ _ acc = return acc
setupBackends (bk : bes) c sb acc = do
  be' ← setupBackend' bk c sb
  setupBackends bes c sb ((bk, be') : acc)
setupBackend' :: BackendKind → Configuration → Switchboard → IO Backend
setupBackend' SwitchboardBK _ _ = error "cannot instantiate a further Switchboard"
setupBackend' EKGViewBK c _ = do
  be :: Cardano.BM.Output ∘ EKGView.EKGView ← Cardano.BM.Output ∘ EKGView.realize c
  return MkBackend

```

```

    {bEffectuate = Cardano.BM.Output ◦ EKGView.effectuate be
    ,bUnrealize = Cardano.BM.Output ◦ EKGView.unrealize be
    }
  setupBackend' AggregationBK c sb = do
    let trace = mainTrace sb
    ctx = TraceContext {loggerName = ""
      ,configuration = c
      ,minSeverity = Debug
      ,tracetype = Neutral
      ,shutdown = pure ()
    }
    be :: Cardano.BM.Output ◦ Aggregation.Aggregation ← Cardano.BM.Output ◦ Aggregation.realizefrom (ctx,
    return MkBackend
      {bEffectuate = Cardano.BM.Output ◦ Aggregation.effectuate be
      ,bUnrealize = Cardano.BM.Output ◦ Aggregation.unrealize be
      }
  setupBackend' KatipBK c _ = do
    be :: Cardano.BM.Output ◦ Log.Log ← Cardano.BM.Output ◦ Log.realize c
    return MkBackend
      {bEffectuate = Cardano.BM.Output ◦ Log.effectuate be
      ,bUnrealize = Cardano.BM.Output ◦ Log.unrealize be
      }

```

1.4.23 Cardano.BM.Output.Log

Internal representation

```

type LogMVar = MVar LogInternal
newtype Log = Log
  {getK :: LogMVar}
data LogInternal = LogInternal
  {kLogEnv :: K.LogEnv
  ,configuration :: Config.Configuration}

```

Log implements effectuate

```

instance IsEffectuator Log where
  effectuate katip item = do
    c ← withMVar (getK katip) $ λk → return (configuration k)
    selscribes ← getScribes c (lnName item)
    forM_ selscribes $ λsc → passN sc katip item

```

Log implements backend functions

```

instance IsBackend Log where
  typeof _ = KatipBK
  realize config = do
    let updateEnv :: K.LogEnv → IO UTCTime → K.LogEnv
        updateEnv le timer =
          le {K._logEnvTimer = timer, K._logEnvHost = "hostname"}
    register :: [ScribeDefinition] → K.LogEnv → IO K.LogEnv
    register [ ] le = return le
    register (defsc : dscs) le = do
      let kind = scKind defsc
          name = scName defsc
          name' = pack (show kind) <> " : " <> name
          scr ← createScribe kind name
          register dscs ≡ K.registerScribe name' scr scribeSettings le
    mockVersion :: Version
    mockVersion = Version [0,1,0,0] [ ]
    scribeSettings :: KC.ScribeSettings
    scribeSettings =
      let bufferSize = 5000 -- size of the queue (in log items)
      in
        KC.ScribeSettings bufferSize
    createScribe FileTextSK name = mkTextFileScribe (FileDescription $ unpack name) False
    createScribe FileJsonSK name = mkJsonFileScribe (FileDescription $ unpack name) False
    createScribe StdoutSK _ = mkStdoutScribe
    createScribe StderrSK _ = mkStderrScribe
    cfoKey ← Config.getOptionOrDefault config (pack "cfokey") (pack "<unknown>")
    le0 ← K.initLogEnv
      (K.Namespace [ "iohk" ])
      (fromString $ (unpack cfoKey) <> " : " <> showVersion mockVersion)
    -- request a new time 'getCurrentTime' at most 100 times a second
    timer ← mkAutoUpdate defaultUpdateSettings {updateAction = getCurrentTime, updateFreq = 10000}
    let le1 = updateEnv le0 timer
    scribes ← getSetupScribes config
    le ← register scribes le1
    kref ← newEmptyMVar
    putMVar kref $ LogInternal le config
    return $ Log kref
  unrealize katip = do
    le ← withMVar (getK katip) $ λk → return (kLogEnv k)
    void $ K.closeScribes le

example :: IO ()
example = do
  config ← Config.setup "from_some_path.yaml"

```

```

k ← setup config
passN (pack (show StdoutSK)) k $ LogNamed
  {lnName = "test"
  ,lnItem = LP $ LogMessage $ LogItem
    {liSelection = Both
    ,liSeverity = Info
    ,liPayload = "Hello!"
    }
  }
passN (pack (show StdoutSK)) k $ LogNamed
  {lnName = "test"
  ,lnItem = LP $ LogValue "cpu-no" 1
  }

-- useful instances for katip
deriving instance K.ToObject LogObject
deriving instance K.ToObject LogItem
deriving instance K.ToObject (Maybe LogObject)
instance KC.LogItem LogObject where
  payloadKeys _ = KC.AllKeys
instance KC.LogItem LogItem where
  payloadKeys _ = KC.AllKeys
instance KC.LogItem (Maybe LogObject) where
  payloadKeys _ = KC.AllKeys

```

Log.passN

The following function copies the **NamedLogItem** to the queues of all scribes that match on their name. Compare start of name of scribe to (*show backend* <> " : "). This function is non-blocking.

```

passN :: Text → Log → NamedLogItem → IO ()
passN backend katip namedLogItem = do
  env ← withMVar (getK katip) $ λk → return (kLogEnv k)
  forM_ (Map.toList $ K._logEnvScribes env) $
    λ(scName, (KC.ScribeHandle _ shChan)) →
      -- check start of name to match ScribeKind
      if backend 'isPrefixOf' scName
      then do
        let item = lnItem namedLogItem
        let (sev,msg,payload) = case item of
          (LP (LogMessage logItem)) →
            (liSeverity logItem, liPayload logItem, Nothing)
          (ObserveDiff counters) →
            let text = toStrict (encodeToLazyText counters)
            in
              (Info, text, Just item)
          (ObserveOpen counters) →

```

```

    let text = toStrict (encodeToLazyText counters)
    in
      (Info, text, Just item)
  (ObserveClose counters) →
    let text = toStrict (encodeToLazyText counters)
    in
      (Info, text, Just item)
  (AggregatedMessage aggregated) →
    let text = T.concat $ (flip map) aggregated $ \ (name, agg) →
      "\n" <> name <> ": " <> pack (show agg)
    in
      (Info, text, Nothing)
  (LP (LogValue name value)) →
    (Debug, name <> " = " <> pack (show value), Nothing)
  KillPill →
    (Info, "Kill pill received!", Nothing)
if (msg ≡ "") ∧ (isNothing payload)
then return ()
else do
  threadIdText ← KC.mkThreadIdText < $ > myThreadId
  let ns = lnName namedLogItem
  itemTime ← env ^. KC.logEnvTimer
  let itemKatip = K.Item {
    _itemApp = env ^. KC.logEnvApp
  , _itemEnv = env ^. KC.logEnvEnv
  , _itemSeverity = sev2klog sev
  , _itemThread = threadIdText
  , _itemHost = env ^. KC.logEnvHost
  , _itemProcess = env ^. KC.logEnvPid
  , _itemPayload = payload
  , _itemMessage = K.logStr msg
  , _itemTime = itemTime
  , _itemNamespace = (env ^. KC.logEnvApp) <> (K.Namespace [ ns ])
  , _itemLoc = Nothing
  }
  void $ atomically $ KC.tryWriteTBQueue shChan (KC.NewItem itemKatip)
else return ()

```

Scribes

```

mkStdoutScribe :: IO K.Scribe
mkStdoutScribe = mkTextFileScribeH stdout True
mkStderrScribe :: IO K.Scribe
mkStderrScribe = mkTextFileScribeH stderr True
mkTextFileScribeH :: Handle → Bool → IO K.Scribe
mkTextFileScribeH handler color = do

```

```

    mkFileScribeH handler formatter color
where
    formatter h colorize verbosity item =
        TIO.hPutStrLn h $! toLazyText $ formatItem colorize verbosity item
mkFileScribeH
    :: Handle
    → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())
    → Bool
    → IO K.Scribe
mkFileScribeH h formatter colorize = do
    hSetBuffering h LineBuffering
    locklocal ← newMVar ()
    let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
        logger item = withMVar locklocal $ \_ →
            formatter h colorize K.V0 item
    pure $ K.Scribe logger (hClose h)
mkTextFileScribe :: FileDescription → Bool → IO K.Scribe
mkTextFileScribe fdesc colorize = do
    mkFileScribe fdesc formatter colorize
where
    formatter :: Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter hdl colorize' v' item =
        case KC._itemMessage item of
            K.LogStr "" →
                -- if message is empty do not output it
                return ()
            _ → do
                let tmsg = toLazyText $ formatItem colorize' v' item
                TIO.hPutStrLn hdl tmsg
mkJsonFileScribe :: FileDescription → Bool → IO K.Scribe
mkJsonFileScribe fdesc colorize = do
    mkFileScribe fdesc formatter colorize
where
    formatter :: (K.LogItem a) ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ()
    formatter h _ verbosity item = do
        let tmsg = case KC._itemMessage item of
            -- if a message is contained in item then only the
            -- message is printed and not the data
            K.LogStr "" → K.itemJson verbosity item
            K.LogStr msg → K.itemJson verbosity $
                item { KC._itemMessage = K.logStr (" " :: Text)
                    , KC._itemPayload = LogItem Both Info $ toStrict $ toLazyText msg
                    }
        TIO.hPutStrLn h (encodeToLazyText tmsg)
mkFileScribe
    :: FileDescription
    → (forall a ◦ K.LogItem a ⇒ Handle → Bool → K.Verbosity → K.Item a → IO ())

```

```

    → Bool
    → IO K.Scribe
mkFileScribe fdesc formatter colorize = do
  let prefixDir = prefixPath fdesc
  (createDirectoryIfMissing True prefixDir)
  'catchIO' (prtoutException ("cannot log prefix directory: " ++ prefixDir))
  let fpath = filePath fdesc
  h ← catchIO (openFile fpath WriteMode) $
    λe → do
      prtoutException ("error while opening log: " ++ fpath) e
      -- fallback to standard output in case of exception
      return stdout
  hSetBuffering h LineBuffering
  scribestate ← newMVar h
  let finalizer :: IO ()
      finalizer = withMVar scribestate hClose
  let logger :: forall a ◦ K.LogItem a ⇒ K.Item a → IO ()
      logger item =
        withMVar scribestate $ λhandler →
          formatter handler colorize K.V0 item
  return $ K.Scribe logger finalizer

formatItem :: Bool → K.Verbosity → K.Item a → Builder
formatItem withColor _verb K.Item {..} =
  fromText header <>
  fromText " " <>
  brackets (fromText timestamp) <>
  fromText " " <>
  KC.unLogStr _itemMessage
where
  header = colorBySeverity _itemSeverity $
    "[" <> mconcat [namedcontext <> ":" <> severity <> ":" <> threadid <> "]"
  namedcontext = KC.intercalateNs _itemNamespace
  severity = KC.renderSeverity _itemSeverity
  threadid = KC.getThreadIdText _itemThread
  timestamp = pack $ formatTime defaultTimeLocale tsformat _itemTime
  tsformat :: String
  tsformat = "%F %T%2Q %Z"
  colorBySeverity s m = case s of
    K.EmergencyS → red m
    K.AlertS     → red m
    K.CriticalS  → red m
    K.ErrorS     → red m
    K.NoticeS    → magenta m
    K.WarningS   → yellow m
    K.InfoS      → blue m
    _            → m
  red = colorize "31"

```



```

    yellow = colorize "33"
    magenta = colorize "35"
    blue = colorize "34"
    colorize c m
      | withColor = "\ESC[" <> c <> "m" <> m <> "\ESC[0m"
      | otherwise = m

-- translate Severity to Log.Severity
sev2klog :: Severity → K.Severity
sev2klog = λcase
  Debug    → K.DebugS
  Info     → K.InfoS
  Notice   → K.NoticeS
  Warning  → K.WarningS
  Error    → K.ErrorS
  Critical → K.CriticalS
  Alert    → K.AlertS
  Emergency → K.EmergencyS

data FileDescription = FileDescription {
  filePath :: !FilePath
  deriving (Show)

  prefixPath :: FileDescription → FilePath
  prefixPath = takeDirectory ∘ filePath

  -- display message and stack trace of exception on stdout
  prtoutException :: Exception e ⇒ String → e → IO ()
  prtoutException msg e = do
    putStrLn msg
    putStrLn ("exception: " ++ displayException e)

```

1.4.24 Cardano.BM.Output.EKGView

Structure of EKGView

```

type EKGViewMVar = MVar EKGViewInternal
newtype EKGView = EKGView
  { getEV :: EKGViewMVar }
data EKGViewInternal = EKGViewInternal
  { evLabels :: HM.HashMap Text Label.Label
  , evServer :: Server
  , evTrace  :: Trace IO
  }

ekgTrace :: EKGView → Configuration → IO (Trace IO)
ekgTrace ekg c = do
  let trace = ekgTrace' ekg

```

```

    ctx = TraceContext {loggerName = " "
                        ,configuration = c
                        ,minSeverity = Debug
                        ,tracetype = Neutral
                        ,shutdown = pure ()
                        }
    Trace.subTrace "#ekgview" (ctx, trace)
where
    ekgTrace' :: EKGView → TraceNamed IO
    ekgTrace' ekgview = BaseTrace.BaseTrace $ Op $ \lognamed → do
        let setlabel :: Text → Text → EKGViewInternal → IO (Maybe EKGViewInternal)
        setlabel name label ekg_i@(EKGViewInternal labels server _) =
            case HM.lookup name labels of
                Nothing → do
                    ekghdl ← getLabel name server
                    Label.set ekghdl label
                    return $ Just $ ekg_i {evLabels = HM.insert name ekghdl labels}
                Just ekghdl → do
                    Label.set ekghdl label
                    return Nothing

        update :: LogObject → LoggerName → EKGViewInternal → IO (Maybe EKGViewInternal)
        update (LP (LogMessage logitem)) logname ekg_i =
            setlabel logname (liPayload logitem) ekg_i
        update (LP (LogValue iname value)) logname ekg_i =
            let logname' = logname <> "." <> iname
            in
                setlabel logname' (pack $ show value) ekg_i
        update _ _ _ = return Nothing

    ekgup ← takeMVar (getEV ekgview)
    let lognam0 = (lnName lognamed)
        lognam1 = case stripPrefix "#ekgview.#aggregation." lognam0 of
            Nothing → lognam0
            Just ln' → ln'
        logname = case stripPrefix "#ekgview." lognam1 of
            Nothing → lognam1
            Just ln' → ln'
    upd ← update (lnItem lognamed) logname ekgup
    case upd of
        Nothing → putMVar (getEV ekgview) ekgup
        Just ekgup' → putMVar (getEV ekgview) ekgup'

```

EKG view is an effectuator

```

instance IsEffectuator EKGView where
    effectuate ekgview item = do
        ekg ← readMVar (getEV ekgview)

```

```

let trace0 = evTrace ekg
trace ← Trace.appendName (lnName item) trace0
case (lnItem item) of
  AggregatedMessage ags → liftIO $ do
    let traceAgg :: [(Text, Aggregated)] → IO ()
    traceAgg [] = return ()
    traceAgg ((n, AggregatedEWMA ewma):r) = do
      trace' ← Trace.appendName n trace
      Trace.traceNamedObject trace' (LP (LogValue "avg" $ avg ewma))
    traceAgg r
  traceAgg ((n, AggregatedStats stats):r) = do
    trace' ← Trace.appendName n trace
    Trace.traceNamedObject trace' (LP (LogValue "mean" (PureD $ meanOfStats stats)))
    Trace.traceNamedObject trace' (LP (LogValue "min" $ fmin stats))
    Trace.traceNamedObject trace' (LP (LogValue "max" $ fmax stats))
    Trace.traceNamedObject trace' (LP (LogValue "count" $ PureI $ fcount stats))
    Trace.traceNamedObject trace' (LP (LogValue "last" $ flast stats))
    Trace.traceNamedObject trace' (LP (LogValue "stdev" (PureD $ stdevOfStats stats)))
    traceAgg r
  traceAgg ags
  _ → liftIO $ Trace.traceNamedObject trace (lnItem item)

```

EKGView implements **Backend** functions

EKGView is an **IsBackend**

```

instance IsBackend EKGView where
  typeof _ = EKGViewBK
  realize config = do
    evref ← newEmptyMVar
    let ekgview = EKGView evref
    evport ← getEKGport config
    ehdl ← forkServer "127.0.0.1" evport
    ekghdl ← getLabel "iohk-monitoring version" ehdl
    Label.set ekghdl $ pack (showVersion version)
    ekgtrace ← ekgTrace ekgview config
    putMVar evref $ EKGViewInternal
      { evLabels = HM.empty
      , evServer = ehdl
      , evTrace = ekgtrace
      }
    return ekgview
  unrealize ekgview = do
    ekg ← takeMVar $ getEV ekgview
    killThread $ serverThreadId $ evServer ekg

```

Interactive testing **EKGView**

```
test :: IO ()
test = do
  c ← Cardano.BM.Configuration.setup "test/config.yaml"
  ev ← Cardano.BM.Output ◦ EKGView.realize c
  effectuate ev $ LogNamed "test.questions" (LP (LogValue "answer" 42))
  effectuate ev $ LogNamed "test.monitor023" (LP (LogMessage (LogItem Public Warning "!!!! ALARM !
```

1.4.25 Cardano.BM.Output.Aggregation

Internal representation

```
type AggregationMVar = MVar AggregationInternal
newtype Aggregation = Aggregation
  {getAg :: AggregationMVar}
data AggregationInternal = AggregationInternal
  {agQueue :: TBQ.TBQueue (Maybe NamedLogItem)
  ,agDispatch :: Async.Async ()
  }
```

Relation from context name to aggregated statistics

We keep the aggregated values (**Aggregated**) for a named context in a *HashMap*.

```
type AggregationMap = HM.HashMap Text AggregatedExpanded
```

Info for Aggregated operations

Apart from the **Aggregated** we keep some valuable info regarding to them; such as when was the last time it was sent.

```
type Timestamp = Word64
data AggregatedExpanded = AggregatedExpanded
  {aeAggregated :: Aggregated
  ,aeResetAfter :: Maybe Integer
  ,aeLastSent :: Timestamp
  }
```

Aggregation implements *effectuate*

Aggregation is an **Accepts a NamedLogItem** Enter the log item into the **Aggregation** queue.

```
instance IsEffectuator Aggregation where
  effectuate agg item = do
    ag ← readMVar (getAg agg)
    atomically $ TBQ.writeTBQueue (agQueue ag) $ Just item
```

Aggregation implements Backend functions**Aggregation is an Declaration of a Backend**

```

instance IsBackend Aggregation where
  typeof _ = AggregationBK
  realize _ = error "Aggregation cannot be instantiated by 'realize'"
  realizefrom trace0@(ctx, _) _ = do
    trace ← Trace.subTrace "#aggregation" trace0
    aggref ← newEmptyMVar
    aggregationQueue ← atomically $ TBQ.newTBQueue 2048
    dispatcher ← spawnDispatcher (configuration ctx) HM.empty aggregationQueue trace
    putMVar aggref $ AggregationInternal aggregationQueue dispatcher
    return $ Aggregation aggref
  unrealize aggregation = do
    let clearMVar :: MVar a → IO ()
      clearMVar = void ∘ tryTakeMVar
    (dispatcher, queue) ← withMVar (getAg aggregation) (λag →
      return (agDispatch ag, agQueue ag))
    -- send terminating item to the queue
    atomically $ TBQ.writeTBQueue queue Nothing
    -- wait for the dispatcher to exit
    res ← Async.waitCatch dispatcher
    either throwM return res
    (clearMVar ∘ getAg) aggregation

```

Asynchronously reading log items from the queue and their processing

```

spawnDispatcher :: Configuration
  → AggregationMap
  → TBQ.TBQueue (Maybe NamedLogItem)
  → Trace.Trace IO
  → IO (Async.Async ())
spawnDispatcher conf aggMap aggregationQueue trace = Async.async $ qProc aggMap
where
  qProc aggregatedMap = do
    maybeItem ← atomically $ TBQ.readTBQueue aggregationQueue
    case maybeItem of
      Just item → do
        (updatedMap, aggregations) ← update (lnItem item) (lnName item) aggregatedMap
        unless (null aggregations) $
          sendAggregated (AggregatedMessage aggregations) (lnName item)
        qProc updatedMap
      Nothing → return ()
  update :: LogObject
    → LoggerName

```

```

→ HM.HashMap Text AggregatedExpanded
→ IO (HM.HashMap Text AggregatedExpanded, [(Text, Aggregated)])
update (LP (LogValue iname value)) logname agmap = do
  let name = logname <> "." <> iname
  aggregated ←
    case HM.lookup name agmap of
      Nothing → do
        -- if Aggregated does not exist; initialize it.
        aggregatedKind ← getAggregatedKind conf name
        case aggregatedKind of
          StatsAK → return $ singletonStats value
          EwmaAK aEWMA → do
            let initEWMA = EmptyEWMA aEWMA
            return $ AggregatedEWMA $ ewma initEWMA value
          Just a → return $ updateAggregation value (aeAggregated a) (aeResetAfter a)
      now ← getMonotonicTimeNSec
  let aggregatedX = AggregatedExpanded {
    aeAggregated = aggregated
    , aeResetAfter = Nothing
    , aeLastSent = now
  }
  namedAggregated = [(iname, aeAggregated aggregatedX)]
  updatedMap = HM.alter (const $ Just $ aggregatedX) name agmap
  -- use of HM.alter so that in future we can clear the Aggregated
  -- by using as alter's arg a function which returns Nothing.
  return (updatedMap, namedAggregated)
update (ObserveDiff counterState) logname agmap = do
  let counters = csCounters counterState
  (mapNew, aggs) ← updateCounters counters logname agmap [ ]
  return (mapNew, reverse aggs)
-- TODO for text messages aggregate on delta of timestamps
update _ _ agmap = return (agmap, [ ])
updateCounters :: [Counter]
→ LoggerName
→ HM.HashMap Text AggregatedExpanded
→ [(Text, Aggregated)]
→ IO (HM.HashMap Text AggregatedExpanded, [(Text, Aggregated)])
updateCounters [ ] _ aggrMap aggs = return $ (aggrMap, aggs)
updateCounters (counter : cs) logname aggrMap aggs = do
  let name = cName counter
  fullname = logname <> "." <> name
  value = cValue counter
  aggregated ←
    case HM.lookup fullname aggrMap of
      -- if Aggregated does not exist; initialize it.
      Nothing → do
        aggregatedKind ← getAggregatedKind conf fullname

```

```

    case aggregatedKind of
      StatsAK → return $ singletonStats value
      EwmaAK aEWMA → do
        let initEWMA = EmptyEWMA aEWMA
        return $ AggregatedEWMA $ ewma initEWMA value
      Just a → return $ updateAggregation value (aeAggregated a) (aeResetAfter a)
now ← getMonotonicTimeNSec
let aggregatedX = AggregatedExpanded {
  aeAggregated = aggregated
  , aeResetAfter = Nothing
  , aeLastSent = now
}
namedAggregated = (((nameCounter counter) <> "." <> name), aggregated)
updatedMap = HM.alter (const $ Just $ aggregatedX) fullname aggrMap
updateCounters cs logname updatedMap (namedAggregated : aggs)
sendAggregated :: LogObject → Text → IO ()
sendAggregated aggregatedMsg@(AggregatedMessage _) logname = do
  -- enter the aggregated message into the Trace
  trace' ← Trace.appendName logname trace
  liftIO $ Trace.traceNamedObject trace' aggregatedMsg
-- ignore every other message that is not of type AggregatedMessage
sendAggregated _ _ = return ()

```

Update aggregation

We distinguish an uninitialized from an already initialized aggregation. The latter is properly initialized.

We use Welford's online algorithm to update the estimation of mean and variance of the sample statistics. (see https://en.wikipedia.org/wiki/Algorithms_for_calculating_variance#Welford's_Online)

```

updateAggregation :: Measurable → Aggregated → Maybe Integer → Aggregated
updateAggregation v (AggregatedStats s) resetAfter =
  let count = fcount s
      reset = maybe False (count ≥) resetAfter
  in
  if reset
  then
    singletonStats v
  else
    let newcount = count + 1
        newvalue = getDouble v
        delta = newvalue - fsum_A s
        dincr = (delta / fromInteger newcount)
        delta2 = newvalue - fsum_A s - dincr
    in
    AggregatedStats Stats {flast = v
      , fmin = min (fmin s) v

```

```

    ,fmax = max (fmax s) v
    ,fcount = newcount
    ,fsum_A = fsum_A s + dincr
    ,fsum_B = fsum_B s + (delta * delta2)
  }
updateAggregation v (AggregatedEWMA e) _ =
  AggregatedEWMA $ ewma e v

```

Calculation of EWMA

Following https://en.wikipedia.org/wiki/Moving_average#Exponential_moving_average we calculate the exponential moving average for a series of values Y_t according to:

$$S_t = \begin{cases} Y_1, & t = 1 \\ \alpha \cdot Y_t + (1 - \alpha) \cdot S_{t-1}, & t > 1 \end{cases}$$

The pattern matching below ensures that the **EWMA** will start with the first value passed in, and will not change type, once determined.

```

ewma :: EWMA → Measurable → EWMA
ewma (EmptyEWMA a) v = EWMA a v
ewma (EWMA a (Microseconds s)) (Microseconds y) =
  EWMA a $ Microseconds $ round $ a * (fromInteger y) + (1 - a) * (fromInteger s)
ewma (EWMA a (Seconds s)) (Seconds y) =
  EWMA a $ Seconds $ round $ a * (fromInteger y) + (1 - a) * (fromInteger s)
ewma (EWMA a (Bytes s)) (Bytes y) =
  EWMA a $ Bytes $ round $ a * (fromInteger y) + (1 - a) * (fromInteger s)
ewma (EWMA a (PureI s)) (PureI y) =
  EWMA a $ PureI $ round $ a * (fromInteger y) + (1 - a) * (fromInteger s)
ewma (EWMA a (PureD s)) (PureD y) =
  EWMA a $ PureD $ a * y + (1 - a) * s
ewma _ _ = error "Cannot average on values of different type"

```


Index

- Aggregated, 31
 - instance of Semigroup, 31
 - instance of Show, 32
- AggregatedExpanded, 58
- Aggregation, 58
- appendName, 14
- Backend, 32
- BaseTrace, 12
 - instance of Contravariant, 13
- Counter, 34
- Counters
 - Dummy
 - readCounters, 22
 - Linux
 - readCounters, 22
- CounterState, 35
- CounterType, 34
- diffCounters, 35
- diffTimeObserved, 20
- evalFilters, 15
- EWMA, 31
- ewma, 62
- getMonoClock, 21
- getOptionOrDefault, 39
- IsBackend, 32
- IsEffectuator, 32
- logAlert, 17
- logAlertP, 17
- logAlertS, 17
- logCritical, 17
- logCriticalP, 17
- logCriticalS, 17
- logDebug, 17
- logDebugP, 17
- logDebugS, 17
- logEmergency, 17
- logEmergencyP, 17
- logEmergencyS, 17
- logError, 17
- logErrorP, 17
- logErrorS, 17
- LoggerName, 36
- logInfo, 17
- logInfoP, 17
- logInfoS, 17
- LogItem, 36
 - liPayload, 36
 - liSelection, 36
 - liSeverity, 36
- LogNamed, 37
- logNotice, 17
- logNoticeP, 17
- logNoticeS, 17
- LogObject, 36
- LogPrims, 36
 - LogMessage, 36
 - LogValue, 36
- LogSelection, 36
 - Both, 36
 - Private, 36
 - Public, 36
 - PublicUnsafe, 36
- logWarning, 17
- logWarningP, 17
- logWarningS, 17
- mainTrace, 47
- Measurable, 28
 - instance of Num, 29
 - instance of Show, 30
- nameCounter, 34
- NamedLogItem, 36
- natTrace, 13
- newContext, 19

- nominalTimeToMicroseconds, 21
- noTrace, 13
- ObservableInstance, 37
 - GhcRtsStats, 37
 - IOStats, 37
 - MemoryStats, 37
 - MonotonicClock, 37
 - ProcessStats, 37
- OutputKind, 37
 - TVarList, 37
 - TVarListNamed, 37
- parseRepresentation, 33
- Port, 33
- readRTSStats, 21
- Representation, 33
- ScribeDefinition, 38
 - scKind, 38
 - scName, 38
 - scRotation, 38
- ScribeId, 37
- ScribeKind
 - FileJsonSK, 37
 - FileTextSK, 37
 - StderrSK, 37
 - StdoutSK, 37
- setupTrace, 18
- Severity, 38
 - Alert, 38
 - Critical, 38
 - Debug, 38
 - Emergency, 38
 - Error, 38
 - Info, 38
 - instance of FromJSON, 38
 - Notice, 38
 - Warning, 38
- singletonStats, 31
- Stats, 30
 - instance of Semigroup, 30
- stats2Text, 31
- stdoutTrace, 15
- SubTrace, 38
 - DropOpening, 38
 - FilterTrace, 38
 - NameOperator, 38
 - NameSelector, 38
 - Neutral, 38
 - NoTrace, 38
 - ObservableTrace, 38
 - TeeTrace, 38
 - UntimedTrace, 38
- subTrace, 18
- Switchboard, 46
 - instance of IsBackend, 47
 - instance of IsEffectuator, 47
 - setupBackends, 48
- Trace, 39
- traceConditionally, 16
- TraceContext, 39
 - configuration, 39
 - loggerName, 39
 - minSeverity, 39
 - shutdown, 39
 - tracetype, 39
- traceInTVar, 16
- traceInTVarIO, 16
- TraceNamed, 39
- traceNamedInTVarIO, 16
- traceNamedItem, 16
- traceNamedObject, 14
- traceWith, 13
- typeofTrace, 13
- updateAggregation, 61
- updateTracetype, 13
- withTrace, 19